

MNIST Classification with Softmax Regression and FCNN

This notebook implements a Softmax regression model and a Fully Connected Neural Network to classify handwritten digits from the MNIST dataset.

Dataset Overview

The MNIST dataset contains 70,000 grayscale images of handwritten digits (0-9), split into 60,000 training and 10,000 test images. Each image is 28x28 pixels. The data is concatenated and split into train:validate:test with ratio 0.6:0.2:0.2 respectively.

Data Loading and Preprocessing

```
import torch
import numpy as np
import pandas as pd
from PIL import Image
import codecs
import matplotlib.pyplot as plt
import math

import kagglehub

# Download latest version
path = kagglehub.dataset_download("hojjatk/mnist-dataset")

print("Path to dataset files:", path)

# copy the downloaded dataset to the current directory
import shutil
shutil.copytree(path, "dataset", dirs_exist_ok=True)

Path to dataset files: /kaggle/input/mnist-dataset
'dataset'
```

Data Reading Utilities

- Custom functions to read MNIST's binary file formats
- `get_int()`: Converts binary data to integers
- `read_label_file()`: Extracts label data
- `read_image_file()`: Extracts image data with proper reshaping

```

def get_int(b):
    return int(codecs.encode(b, 'hex'), 16)

def read_label_file(path):
    with open(path, 'rb') as f:
        data = f.read()
        assert get_int(data[:4]) == 2049
        length = get_int(data[4:8])
        parsed = np.frombuffer(data, dtype=np.uint8, offset=8)
        return np.array(parsed, dtype=np.uint8).reshape(length)

def read_image_file(path):
    with open(path, 'rb') as f:
        data = f.read()
        assert get_int(data[:4]) == 2051
        length = get_int(data[4:8])
        num_rows = get_int(data[8:12])
        num_cols = get_int(data[12:16])
        parsed = np.frombuffer(data, dtype=np.uint8, offset=16)
        return np.array(parsed,
dtype=np.uint8).reshape((length,num_rows,num_cols))

```

Data Loading

```

train_images = read_image_file('dataset/train-images.idx3-ubyte')
train_labels = read_label_file('dataset/train-labels.idx1-ubyte')
test_images = read_image_file('dataset/t10k-images.idx3-ubyte')
test_labels = read_label_file('dataset/t10k-labels.idx1-ubyte')

print(train_images.shape, train_labels.shape)

(60000, 28, 28) (60000,)

# concat train and test
images = np.concatenate([train_images, test_images])
labels = np.concatenate([train_labels, test_labels])

```

Train-Validation-Test Split

- Uses scikit-learn's train_test_split
- Creates training set, validation set, test set

```

from sklearn.model_selection import train_test_split
train_test_images, val_images, train_test_labels, val_labels =
train_test_split(images, labels, test_size=0.2, random_state=42)
train_images, test_images, train_labels, test_labels =
train_test_split(train_test_images, train_test_labels, test_size=0.2,
random_state=42)

print(train_images.shape, train_labels.shape)

```

```
(44800, 28, 28) (44800,)
```

Data Preprocessing

Custom Dataset Class

- Create a PyTorch Dataset class for MNIST
- Implement `__init__`, `__getitem__`, and `__len__` methods
- Add image transformations

```
from torch.utils.data import DataLoader, Dataset
from torchvision.transforms import Compose, ToTensor, Normalize

class MNISTDataset(Dataset):
    def __init__(self, images, labels, image_transform:Compose =
None):
        self.images = images
        self.labels = labels
        self.image_transform = image_transform

    def __getitem__(self, index: int):
        image, label = self.images[index], self.labels[index]
        image = Image.fromarray(image, mode='L')
        if self.image_transform is not None:
            image = self.image_transform(image)

        return image, label

    def __len__(self):
        return len(self.images)
```

Image Transformations

- Define transformation pipelines using Compose
- Include:
 - ToTensor conversion
 - Normalization
 - Optional augmentations for training RandomAffine, GaussianBlur, etc.

```
image_transform = Compose([
    ToTensor(),
    Normalize((0.1307,), (0.3081,))
])

from torchvision.transforms import (
    RandomAffine,
    GaussianBlur,
    ElasticTransform
)
```

```

image_augmentation_transform = Compose([
    ToTensor(),
    RandomAffine(
        degrees=7.5,
        translate=(0.1, 0.1),
        scale=(0.9, 1.1),
        shear=5,
    ),
    GaussianBlur(kernel_size=3, sigma=(0.1, 0.2)),
    ElasticTransform(alpha=37.0, sigma=5.0),
    Normalize((0.1307,), (0.3081,))
])

train_dataset = MNISTDataset(train_images, train_labels,
image_transform)
train_augmented_dataset = MNISTDataset(train_images, train_labels,
image_augmentation_transform)
val_dataset = MNISTDataset(val_images, val_labels, image_transform)
test_dataset = MNISTDataset(test_images, test_labels, image_transform)

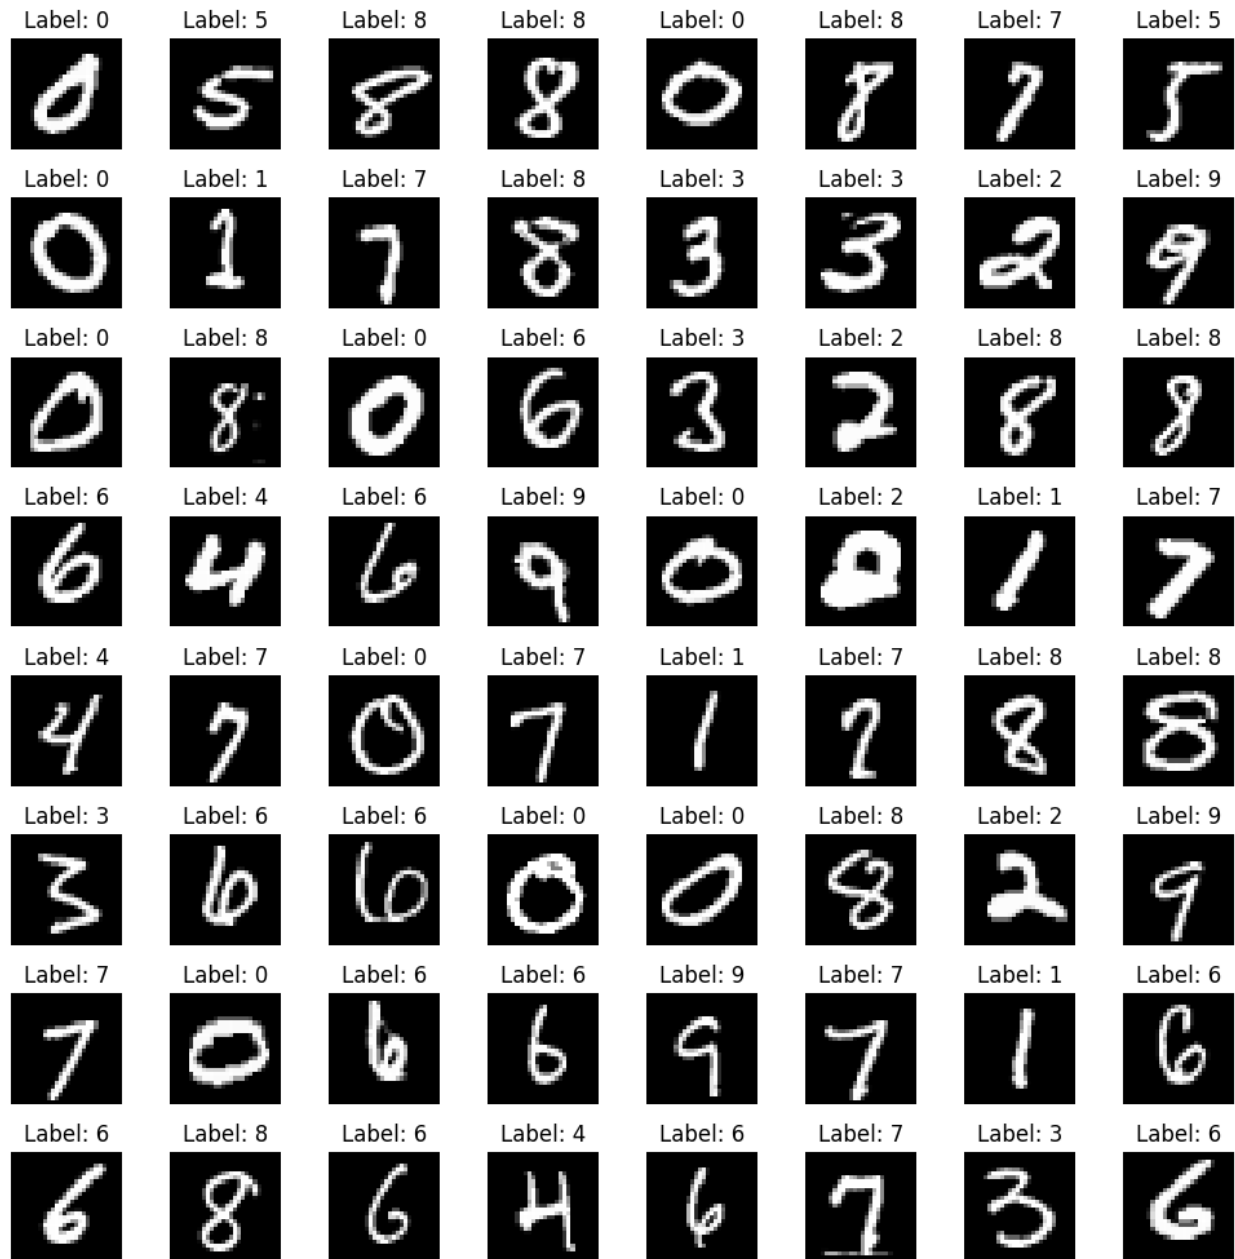
train_loader = DataLoader(train_dataset, batch_size=64, num_workers=4,
prefetch_factor=2, shuffle=True)
train_augmented_loader = DataLoader(train_augmented_dataset,
batch_size=64, num_workers=4, prefetch_factor=2, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=64, num_workers=4,
prefetch_factor=2)
test_loader = DataLoader(test_dataset, batch_size=64, num_workers=4,
prefetch_factor=2)

def visualize_batch(loader):
    batch_images, batch_labels = next(iter(loader))
    print(batch_images.size(), batch_labels.size())
    batch_size = batch_images.size(0)
    cols = math.ceil(math.sqrt(batch_size))
    rows = math.ceil(batch_size / cols)
    # Create a grid of images
    fig, axes = plt.subplots(rows, cols, figsize=(10, 10))
    for i, (image, label) in enumerate(zip(batch_images,
batch_labels)):
        ax = axes[i//rows, i%cols]
        # Convert tensor back to image for display
        ax.imshow(image.squeeze(), cmap='gray')
        ax.axis('off')
        ax.set_title(f'Label: {label.item()}')
    plt.tight_layout()
    plt.show()

visualize_batch(train_loader)

torch.Size([64, 1, 28, 28]) torch.Size([64])

```



Training

Training Configuration

- Create a dataclass for experiment configuration
- Include hyperparameters:
 - Learning rates
 - Batch sizes
 - Dropout rates

- Optimization algorithms
- Scheduling strategies

```

from dataclasses import dataclass, field
from typing import List

@dataclass
class TrainingConfig:
    # Model parameters
    input_size: int = 28 * 28
    output_size: int = 10
    hidden_sizes: List[int] = field(default_factory=lambda: []) #
    Empty by default for softmax
    dropout_rate: float = 0.0

    # Training parameters
    batch_size: int = 128
    num_epochs: int = 50
    learning_rate: float = 0.01
    weight_decay: float = 0.0

    # Optimizer parameters
    optimizer: str = "sgd" # "sgd", "adam"

    # LR scheduler parameters
    scheduler_type: str = "none" # "none", "cosine", "warmup_cosine"
    warmup_epochs: int = 5
    min_lr: float = 1e-5

    # Early stopping parameters
    early_stopping: bool = True
    patience: int = 10
    loss_delta: float = 0.02

    # DataLoader parameters
    num_workers: int = 4
    prefetch_factor: int = 2
    use_augmentation: bool = False

    # Device
    device: str = "cuda" if torch.cuda.is_available() else "cpu"

    # Experiment name
    experiment_name: str = "mnist_exp"
    project_name: str = "mnist-exp"

```

Model Architectures

- Implement two model types:
 - Softmax Regression

- Simple linear model
- No hidden layers
- Deep Neural Network based on [Deep Big Simple Neural Nets Excel on Handwritten Digit Recognition \(Dan Claudiu Ciresan, et al.\)](#)
 - Multiple configurable hidden layers
 - Batch normalization
 - Dropout
 - Kaiming He initialization
 - ReLU activation

```
class SoftmaxRegression(torch.nn.Module):
    def __init__(self, input_size, output_size, dropout_rate=0.0):
        super(SoftmaxRegression, self).__init__()
        self.input_size = input_size
        self.output_size = output_size
        self.dropout = torch.nn.Dropout(dropout_rate)
        self.linear = torch.nn.Linear(input_size, output_size)

    def forward(self, x):
        x = x.view(-1, self.input_size) # remove the channel dimension, flatten the image
        x = self.dropout(x)
        return self.linear(x)

class DeepBigSimpleNet(torch.nn.Module):
    def __init__(self, input_size=784, hidden_sizes=[2500, 2000, 1500], num_classes=10, dropout_rate=0.5):
        super(DeepBigSimpleNet, self).__init__()
        self.input_size = input_size
        self.hidden_sizes = hidden_sizes
        self.num_classes = num_classes

        # Create layers list
        layers = []

        # Input layer
        layers.append(torch.nn.Linear(input_size, hidden_sizes[0]))
        layers.append(torch.nn.BatchNorm1d(hidden_sizes[0]))
        layers.append(torch.nn.ReLU())
        layers.append(torch.nn.Dropout(dropout_rate))

        # Hidden layers with decreasing sizes
        for i in range(len(hidden_sizes) - 1):
            layers.append(torch.nn.Linear(hidden_sizes[i], hidden_sizes[i+1]))
            layers.append(torch.nn.BatchNorm1d(hidden_sizes[i+1]))
            layers.append(torch.nn.ReLU())
            layers.append(torch.nn.Dropout(dropout_rate))
```

```

    # Output layer
    layers.append(torch.nn.Linear(hidden_sizes[-1], num_classes))

    self.model = torch.nn.Sequential(*layers)
    self._initialize_weights()

    def _initialize_weights(self):
        for m in self.modules():
            if isinstance(m, torch.nn.Linear):
                torch.nn.init.kaiming_normal_(m.weight, mode='fan_in',
nonlinearity='relu')
                if m.bias is not None:
                    torch.nn.init.constant_(m.bias, 0)
            elif isinstance(m, torch.nn.BatchNorm1d):
                torch.nn.init.constant_(m.weight, 1)
                torch.nn.init.constant_(m.bias, 0)

    def forward(self, x):
        x = x.view(-1, self.input_size) # Flatten the input
        return self.model(x)

from torch.nn import CrossEntropyLoss
import torch.nn.functional as F
import torch.optim as optim
from torch.optim.lr_scheduler import LinearLR, CosineAnnealingLR,
SequentialLR
from tqdm import tqdm
from tabulate import tabulate
import wandb
from time import time

```

Logging and Visualization

1. [Wandb integration for experiment tracking](#)
 - Log and visualize:
 - Hyperparameters
 - Metrics
 - Visualizations
 - Model performance
 - Misclassifications
 - Confusion matrix
 - Gradient and weights histograms
2. Create experiment tracker
 - Log and visualize:
 - Training curves
 - Loss progression
 - Accuracy metrics

- Confusion matrices

```
wandb.login(key="d8ff0fac98c036a4ac0587814c4fd1a2e60f2512")

wandb: Using wandb-core as the SDK backend. Please refer to
https://wandb.me/wandb-core for more information.
wandb: Currently logged in as: ahmedayman4a77. Use `wandb login --
relogin` to force relogin
wandb: WARNING If you're specifying your api key in code, ensure this
code is not shared publicly.
wandb: WARNING Consider setting the WANDB_API_KEY environment
variable, or running `wandb login` from the command line.
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
```

True

```
import wandb
import torch
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix

class WandBLogger:
    def __init__(self, config):
        self.run = wandb.init(
            project=config.project_name,
            name=config.experiment_name,
            config=config.__dict__,
        )

    def log_model_graph(self, model, loss):
        wandb.watch(model, criterion=loss, log="all", log_freq=100)

    def log_batch_predictions(self, images, labels, predictions,
step):
        # Log sample predictions
        fig, ax = plt.subplots(4, 4, figsize=(10, 10))
        for idx in range(16):
            i, j = idx//4, idx%4
            ax[i,j].imshow(images[idx].squeeze(), cmap='gray')
            ax[i,j].set_title(f'Pred: {predictions[idx]}\nTrue:
{labels[idx]}')
            ax[i,j].axis('off')
        wandb.log({"predictions": wandb.Image(fig)}, commit=False)
        plt.close()

    def log_confusion_matrix(self, true_labels, predictions, step):
        cm = confusion_matrix(true_labels, predictions)
        fig = plt.figure(figsize=(10, 10))
        sns.heatmap(cm, annot=True, fmt='d')
        wandb.log({"confusion_matrix": wandb.Image(fig)},
```

```

commit=False)
    plt.close()

    def log_grad_flow(self, named_parameters):
        ave_grads = []
        layers = []
        for n, p in named_parameters:
            if p.requires_grad and p.grad is not None:
                layers.append(n)
                ave_grads.append(p.grad.abs().mean().item())
        fig = plt.figure(figsize=(10, 5))
        plt.plot(ave_grads, alpha=0.3, color="b")
        plt.title("Gradient Flow")
        wandb.log({"grad_flow": wandb.Image(fig)}, commit=False)
        plt.close()

    def log_misclassified(self, images, labels, predictions, step):
        mask = predictions != labels
        if not mask.any():
            return

        misclassified_images = images[mask][:16]
        misclassified_labels = labels[mask][:16]
        misclassified_preds = predictions[mask][:16]

        fig, ax = plt.subplots(4, 4, figsize=(10, 10))
        for idx in range(min(16, len(misclassified_images))):
            i, j = idx//4, idx%4
            ax[i,j].imshow(misclassified_images[idx].squeeze(),
cmap='gray')
            ax[i,j].set_title(f'Pred: {misclassified_preds[idx]}\
nTrue: {misclassified_labels[idx]}')
            ax[i,j].axis('off')
            wandb.log({"misclassified": wandb.Image(fig)}, commit=False)
        plt.close()

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from typing import Dict, List

class ExperimentTracker:
    def __init__(self):
        self.experiment_metrics = {} # Store per-epoch metrics
        self.experiment_summaries = [] # Store final results

    def log_epoch(self, experiment_name: str, epoch_metrics: Dict):
        if experiment_name not in self.experiment_metrics:
            self.experiment_metrics[experiment_name] = []
        self.experiment_metrics[experiment_name].append(epoch_metrics)

```

```

def log_summary(self, experiment_name: str, summary_metrics:
Dict):
    summary_metrics['experiment_name'] = experiment_name
    self.experiment_summaries.append(summary_metrics)

def get_experiment_df(self, experiment_name: str) -> pd.DataFrame:
    return pd.DataFrame(self.experiment_metrics[experiment_name])

def get_all_summaries(self) -> pd.DataFrame:
    return pd.DataFrame(self.experiment_summaries)

def plot_training_curves(self, metric: str = 'loss'):
    plt.figure(figsize=(12, 6))
    for exp_name, metrics in self.experiment_metrics.items():
        df = pd.DataFrame(metrics)
        plt.plot(df['epoch'], df[f'train_{metric}'],
label=f'{exp_name} (train)')
        plt.plot(df['epoch'], df[f'val_{metric}'],
label=f'{exp_name} (val)')
        plt.xlabel('Epoch')
        plt.ylabel(metric.capitalize())
        plt.title(f'Training and Validation {metric.capitalize()}
Curves')
    plt.legend()
    plt.grid(True)
    return plt.gcf()

def plot_confusion_matrix(self, experiment_name: str):
    summary_df = self.get_all_summaries()
    cm = summary_df[summary_df['experiment_name'] ==
experiment_name]['confusion_matrix'].values[0]
    plt.figure(figsize=(10, 10))
    sns.heatmap(cm, annot=True, fmt='d')
    plt.title(f'Confusion Matrix for {experiment_name}')
    return plt.gcf()

def plot_experiment_comparison(self, metric: str):
    summary_df = self.get_all_summaries()
    plt.figure(figsize=(10, 6))
    sns.barplot(data=summary_df, x='experiment_name', y=metric)
    plt.xticks(rotation=45)
    plt.title(f'Comparison of {metric} across experiments')
    return plt.gcf()

```

Early Stopping

- Monitor validation loss
- Stop training if loss does not improve after a certain number of epochs

```

class EarlyStopper:
    def __init__(self, patience=1, min_delta=0):
        self.patience = patience
        self.min_delta = min_delta
        self.counter = 0
        self.min_validation_loss = float('inf')

    def __call__(self, validation_loss):
        if validation_loss < self.min_validation_loss:
            self.min_validation_loss = validation_loss
            self.counter = 0
        elif validation_loss > (self.min_validation_loss +
self.min_delta):
            self.counter += 1
            if self.counter >= self.patience:
                return True
            return False

```

Evaluator

Track:

- Accuracy
- Inference time
- Confusion matrix

```

from sklearn.metrics import precision_recall_fscore_support,
confusion_matrix

class Evaluator:
    def __init__(self, model, test_loader, device):
        self.model = model
        self.test_loader = test_loader
        self.device = device

    def __call__(self):
        self.model.eval()
        all_preds = []
        all_labels = []
        inference_times = []

        with torch.no_grad():
            for images, labels in self.test_loader:
                images = images.to(self.device)
                start_time = time()
                outputs = self.model(images)
                inference_times.append(time() - start_time)

                preds = outputs.argmax(dim=1)

```

```

        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.numpy())

    metrics = self.compute_metrics(all_labels, all_preds,
inference_times)
    return metrics

    def compute_metrics(self, labels, preds, times):

        return {
            'accuracy': (np.array(labels) == np.array(preds)).mean(),
            'confusion_matrix': confusion_matrix(labels, preds),
            'avg_inference_time': np.mean(times),
            'std_inference_time': np.std(times)
        }

```

Training Infrastructure

Develop training manager class that includes:

- Optimizer selection
- Learning rate scheduling
- Early stopping
- Logging mechanisms
- Training loop
- Validation loop
- Wandb integration

```

class TrainingManager:
    def __init__(self, config: TrainingConfig, tracker:
ExperimentTracker):
        self.config = config
        self.logger = WandBLogger(config)
        self.tracker = tracker
        if self.config.early_stopping:
            self.early_stopper =
EarlyStopper(patience=self.config.patience,
min_delta=self.config.loss_delta)
        else:
            self.early_stopper = None

    def get_scheduler(self, optimizer):
        if self.config.scheduler_type == "none":
            return None
        elif self.config.scheduler_type == "cosine":
            return torch.optim.lr_scheduler.CosineAnnealingLR(
optimizer,
T_max=self.config.num_epochs,

```

```

        eta_min=self.config.min_lr
    )
    elif self.config.scheduler_type == "warmup_cosine":
        warmup_scheduler = LinearLR(optimizer,
total_iters=self.config.warmup_epochs)
        main_scheduler = CosineAnnealingLR(optimizer,
T_max=self.config.num_epochs - self.config.warmup_epochs,
eta_min=self.config.min_lr)
        scheduler = SequentialLR(optimizer,
schedulers=[warmup_scheduler, main_scheduler],
milestones=[self.config.warmup_epochs])
        return scheduler

def train(self, model, train_loader, val_loader, test_loader):
    model = model.to(self.config.device)
    if self.config.optimizer == "sgd":
        optimizer = torch.optim.SGD(
            model.parameters(),
            lr=self.config.learning_rate,
            weight_decay=self.config.weight_decay
        )
    elif self.config.optimizer == "adam":
        optimizer = torch.optim.Adam(
            model.parameters(),
            lr=self.config.learning_rate,
            weight_decay=self.config.weight_decay
        )
    else:
        raise ValueError("Invalid optimizer type")

    scheduler = self.get_scheduler(optimizer)
    criterion = CrossEntropyLoss()
    self.logger.log_model_graph(model, criterion)

    start_time = time()
    for epoch in range(self.config.num_epochs):
        epoch_start = time()
        # Training
        model.train()
        train_loss = 0
        correct = 0
        total = 0

        for batch_idx, (inputs, targets) in
enumerate(train_loader):
            inputs, targets = inputs.to(self.config.device),
targets.to(self.config.device)

```

```

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()

    train_acc = 100. * correct / total
    train_loss = train_loss / len(train_loader)

    # Validation
    val_loss, val_acc = self.validate(model, val_loader,
criterion, epoch)

    # Update scheduler
    if scheduler is not None:
        scheduler.step()

    epoch_time = time() - epoch_start

    epoch_metrics = {
        "epoch": epoch,
        "train_loss": train_loss,
        "train_acc": train_acc,
        "val_loss": val_loss,
        "val_acc": val_acc,
        "learning_rate": optimizer.param_groups[0]['lr'],
        'epoch_time': epoch_time,
    }
    # Log metrics
    self.tracker.log_epoch(self.config.experiment_name,
epoch_metrics)
    wandb.log(epoch_metrics)

    if self.early_stopper is not None and
self.early_stopper(val_loss):
        print(f"Early stopping triggered at epoch {epoch}")
        break

    total_time = time() - start_time

    # Final evaluation
    evaluate = Evaluator(model, test_loader, self.config.device)
    test_metrics = evaluate()

```

```

        test_metrics['total_time'] = total_time
        self.tracker.log_summary(self.config.experiment_name,
test_metrics)

        fig = plt.figure(figsize=(10, 10))
        sns.heatmap(test_metrics['confusion_matrix'], annot=True,
fmt='d')
        # Log final metrics
        wandb.log({
            'total_training_time': total_time,
            'test_accuracy': test_metrics['accuracy'],
            'avg_inference_time': test_metrics['avg_inference_time'],
            'test_confusion_matrix': wandb.Image(fig)
        })

        wandb.finish()
        return test_metrics

def validate(self, model, val_loader, criterion, epoch):
    model.eval()
    val_loss = 0
    correct = 0
    total = 0
    val_images, val_labels = [], []
    val_preds = []

    with torch.no_grad():
        for inputs, targets in val_loader:
            inputs, targets = inputs.to(self.config.device),
targets.to(self.config.device)
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            val_loss += loss.item()
            predictions = outputs.argmax(dim=1)
            total += targets.size(0)
            correct += predictions.eq(targets).sum().item()
            val_images.extend(inputs.cpu())
            val_labels.extend(targets.cpu())
            val_preds.extend(predictions.cpu())

    val_images = torch.stack(val_images)
    val_labels = torch.tensor(val_labels)
    val_preds = torch.tensor(val_preds)

    # Log validation artifacts
    self.logger.log_batch_predictions(
        val_images[:16],

```



```

        val_labels[:16],
        val_preds[:16],
        epoch
    )
    self.logger.log_confusion_matrix(
        val_labels,
        val_preds,
        epoch
    )
    self.logger.log_misclassified(
        val_images,
        val_labels,
        val_preds,
        epoch
    )
    self.logger.log_grad_flow(model.named_parameters())

    return val_loss / len(val_loader), 100. * correct / total

```

Experiment Runner

Create function to run multiple experiments with different configurations of:

- Model architectures
- Hyperparameters
- Optimization strategies
- Scheduling strategies
- Regularization

```

experiments = [
    TrainingConfig(learning_rate=1.0,
experiment_name="softmax_lr_1.0", early_stopping=False),
    TrainingConfig(learning_rate=0.01,
experiment_name="softmax_lr_0.01"),
    TrainingConfig(learning_rate=0.00001,
experiment_name="softmax_lr_0.00001"),
    TrainingConfig(
        learning_rate=0.01,
        scheduler_type="warmup_cosine",
        warmup_epochs=5,
        experiment_name="softmax_warmup_cosine"
    ),
    TrainingConfig(batch_size=1, experiment_name="softmax_batch_1"),
    TrainingConfig(batch_size=32, experiment_name="softmax_batch_32"),
    TrainingConfig(batch_size=256,
experiment_name="softmax_batch_256"),
    TrainingConfig(
        learning_rate=0.01,
        weight_decay=0.01,

```

```

        dropout_rate=0.1,
        batch_size=256,
        experiment_name="softmax_with_regularization"
    ),
    # DeepBigSimpleNet experiments
    TrainingConfig(
        input_size=28 * 28,
        output_size=10,
        hidden_sizes=[40, 15],
        dropout_rate=0.1,
        learning_rate=0.01,
        weight_decay=1e-4,
        batch_size=128,
        scheduler_type="warmup_cosine",
        experiment_name="deep_net_nano"
    ),
    TrainingConfig(
        input_size=28 * 28,
        output_size=10,
        hidden_sizes=[1500, 1000, 500],
        dropout_rate=0.1,
        learning_rate=0.01,
        weight_decay=1e-4,
        batch_size=128,
        scheduler_type="warmup_cosine",
        optimizer="adam",
        use_augmentation=True,
        experiment_name="deep_net_medium"
    ),
    TrainingConfig(
        input_size=28 * 28,
        output_size=10,
        hidden_sizes=[2500, 2000, 1500, 1000, 500],
        dropout_rate=0.1,
        learning_rate=0.01,
        weight_decay=1e-4,
        batch_size=128,
        scheduler_type="warmup_cosine",
        optimizer="adam",
        use_augmentation=True,
        experiment_name="deep_net_large"
    )
]

def run_experiments(experiments, tracker=None):
    if tracker is None:
        tracker = ExperimentTracker()

    for config in experiments:
        print(f"Running experiment: {config.experiment_name}")

```

```

        chose_train_dataset = train_augmented_dataset if
config.use_augmentation else train_dataset
    # Create data loaders
    train_loader = DataLoader(
        chose_train_dataset,
        batch_size=config.batch_size,
        num_workers=config.num_workers,
        prefetch_factor=config.prefetch_factor,
        shuffle=True
    )
    val_loader = DataLoader(
        val_dataset,
        batch_size=config.batch_size,
        num_workers=config.num_workers,
        prefetch_factor=config.prefetch_factor
    )
    test_loader = DataLoader(
        test_dataset,
        batch_size=config.batch_size,
        num_workers=config.num_workers,
        prefetch_factor=config.prefetch_factor
    )

    # Create appropriate model based on config
    if len(config.hidden_sizes) > 0:
        # Use DeepBigSimpleNet if hidden_sizes are specified
        model = DeepBigSimpleNet(
            input_size=config.input_size,
            hidden_sizes=config.hidden_sizes,
            num_classes=config.output_size,
            dropout_rate=config.dropout_rate
        )
    else:
        # Use SoftmaxRegression for basic experiments
        model = SoftmaxRegression(
            input_size=config.input_size,
            output_size=config.output_size,
            dropout_rate=config.dropout_rate
        )

    trainer = TrainingManager(config, tracker)
    trainer.train(model, train_loader, val_loader, test_loader)

```

Running Experiments

All experiments are run using the experiment runner function. The results are logged and visualized using Wandb. You can view the results [here](#).

```
tracker = ExperimentTracker()  
run_experiments(experiments, tracker)
```

Running experiment: softmax_lr_1.0

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Running experiment: softmax_lr_0.01

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Running experiment: softmax_lr_0.00001

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: softmax_warmup_cosine
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

/usr/local/lib/python3.10/dist-packages/torch/optim/
lr_scheduler.py:232: UserWarning: The epoch parameter in
`scheduler.step()` was not necessary and is being deprecated where
possible. Please use `scheduler.step()` to step the scheduler. During
the deprecation, if epoch is different from None, the closed form is
used instead of the new chainable form, where available. Please open
an issue if you are unable to replicate your use case:
https://github.com/pytorch/pytorch/issues/new/choose.
  warnings.warn(EPOCH_DEPRECATION_WARNING, UserWarning)

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: softmax_batch_1
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Early stopping triggered at epoch 10
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: softmax_batch_32
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: softmax_batch_256
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: softmax_with_regularization
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: deep_net_nano
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

/usr/local/lib/python3.10/dist-packages/torch/optim/
lr_scheduler.py:232: UserWarning: The epoch parameter in
`scheduler.step()` was not necessary and is being deprecated where
possible. Please use `scheduler.step()` to step the scheduler. During
the deprecation, if epoch is different from None, the closed form is
used instead of the new chainable form, where available. Please open
an issue if you are unable to replicate your use case:
https://github.com/pytorch/pytorch/issues/new/choose.
  warnings.warn(EPOCH_DEPRECATION_WARNING, UserWarning)

<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
Running experiment: deep_net_medium
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>
<IPython.core.display.HTML object>

/usr/local/lib/python3.10/dist-packages/torch/optim/
lr_scheduler.py:232: UserWarning: The epoch parameter in
```

```
`scheduler.step()` was not necessary and is being deprecated where possible. Please use `scheduler.step()` to step the scheduler. During the deprecation, if epoch is different from None, the closed form is used instead of the new chainable form, where available. Please open an issue if you are unable to replicate your use case:  
https://github.com/pytorch/pytorch/issues/new/choose.
```

```
warnings.warn(EPOCH_DEPRECATION_WARNING, UserWarning)
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
Running experiment: deep_net_large
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
/usr/local/lib/python3.10/dist-packages/torch/optim/  
lr_scheduler.py:232: UserWarning: The epoch parameter in  
`scheduler.step()` was not necessary and is being deprecated where possible. Please use `scheduler.step()` to step the scheduler. During the deprecation, if epoch is different from None, the closed form is used instead of the new chainable form, where available. Please open an issue if you are unable to replicate your use case:  
https://github.com/pytorch/pytorch/issues/new/choose.
```

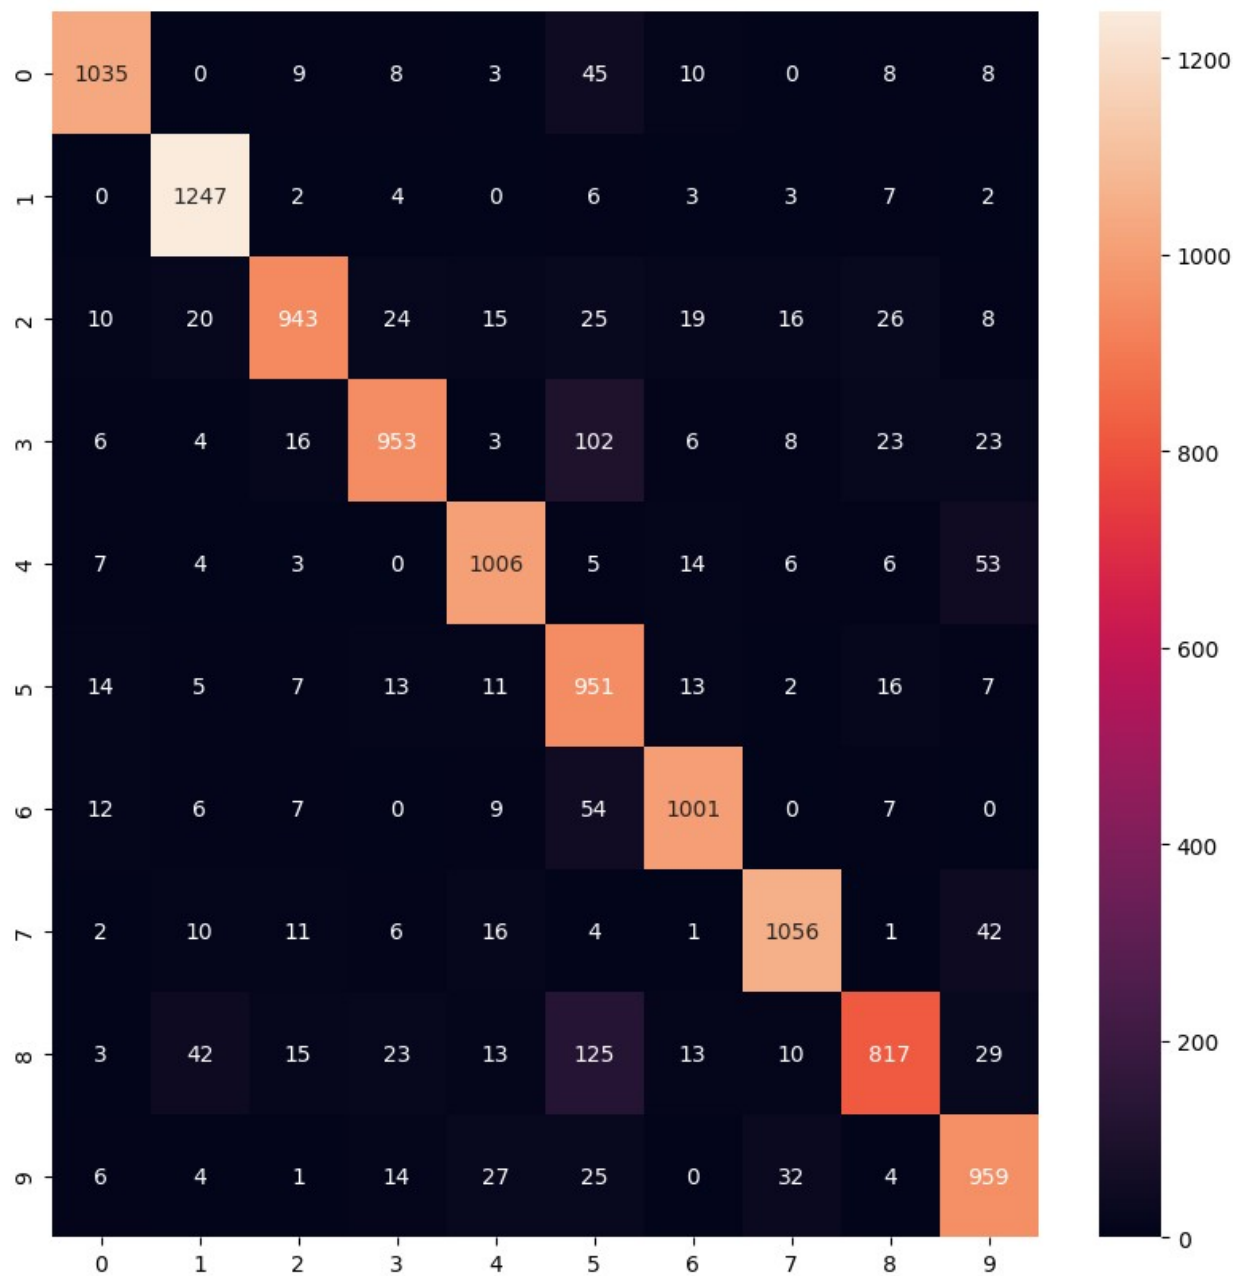
```
warnings.warn(EPOCH_DEPRECATION_WARNING, UserWarning)
```

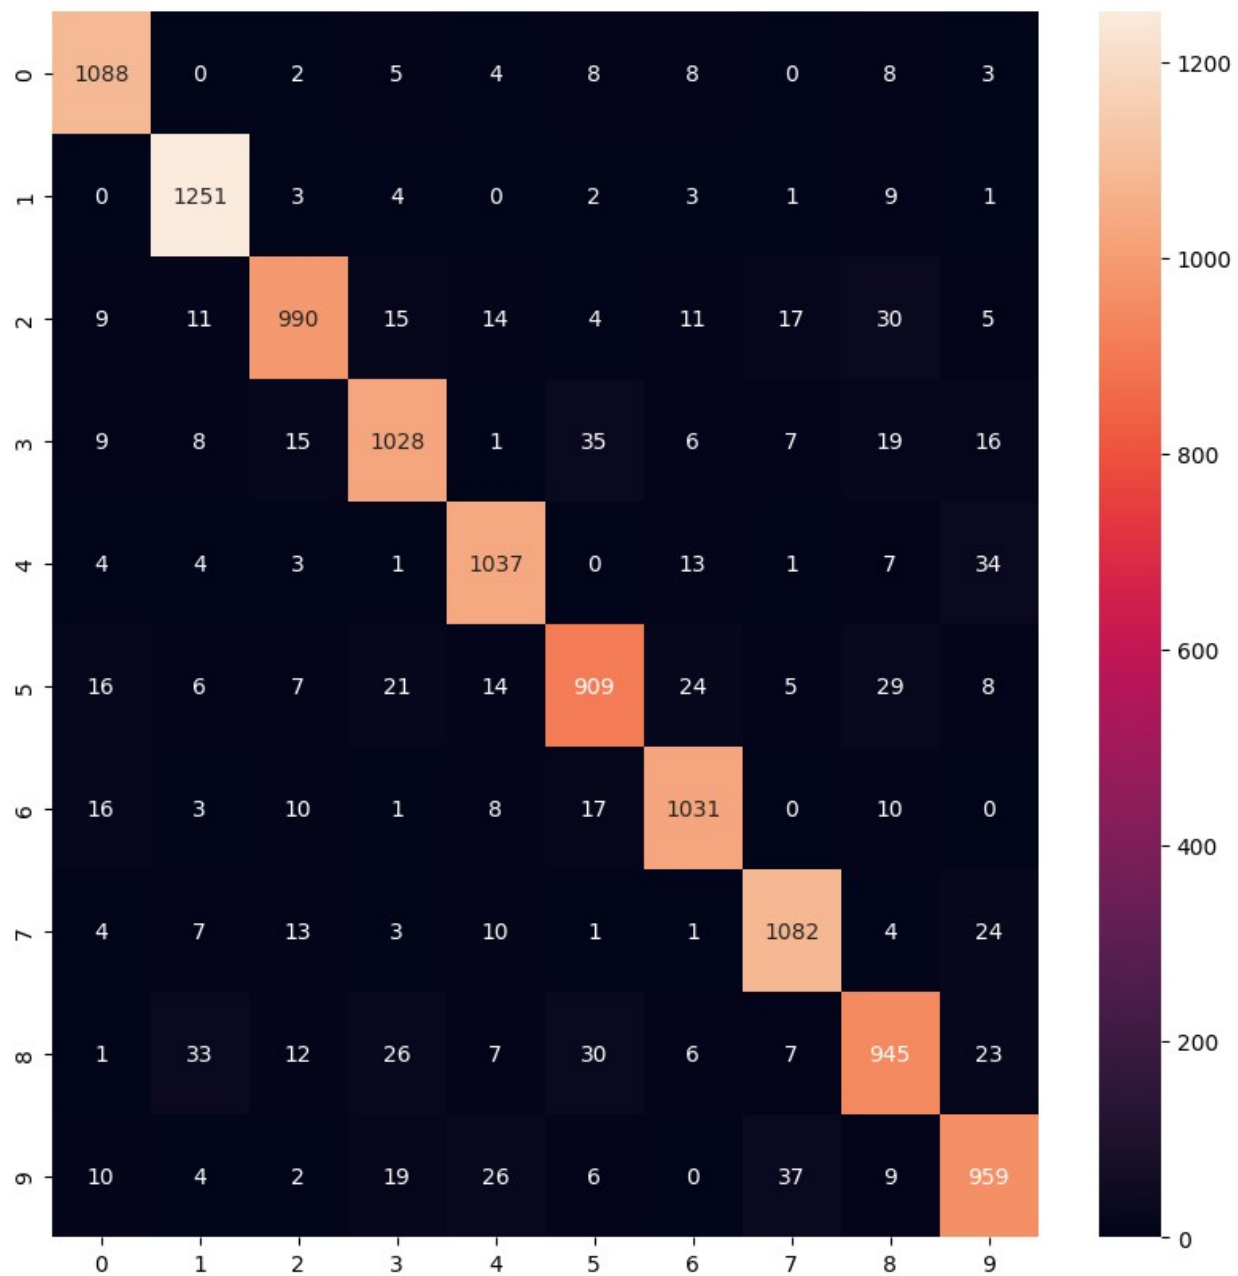
```
<IPython.core.display.HTML object>
```

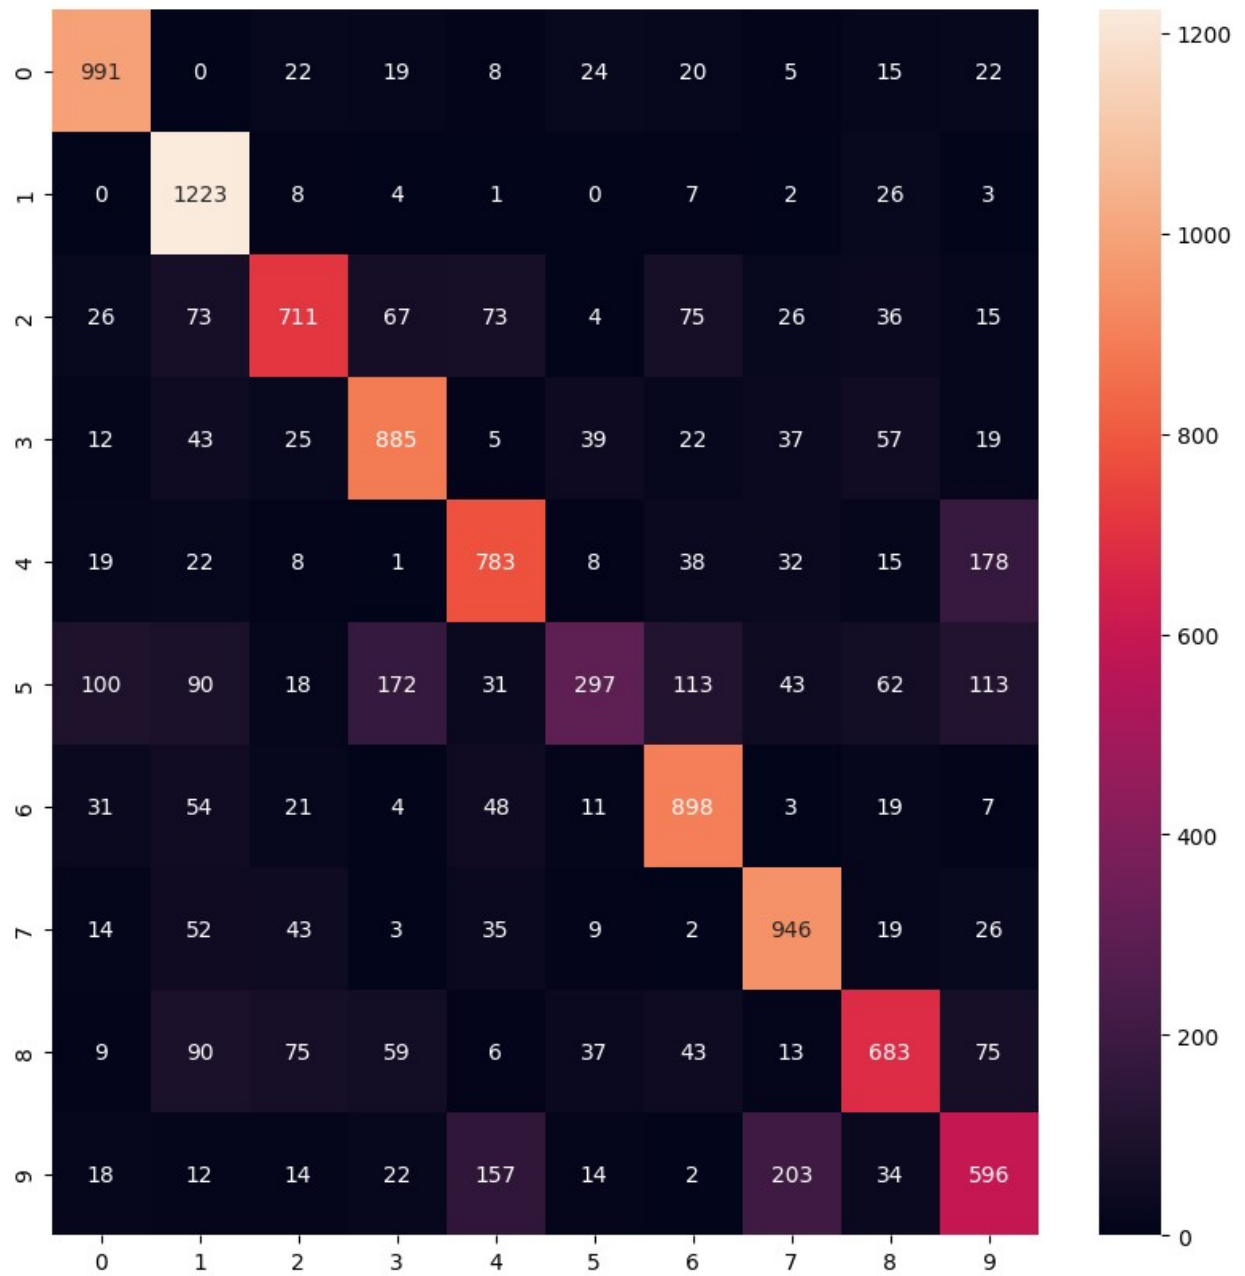
```
<IPython.core.display.HTML object>
```

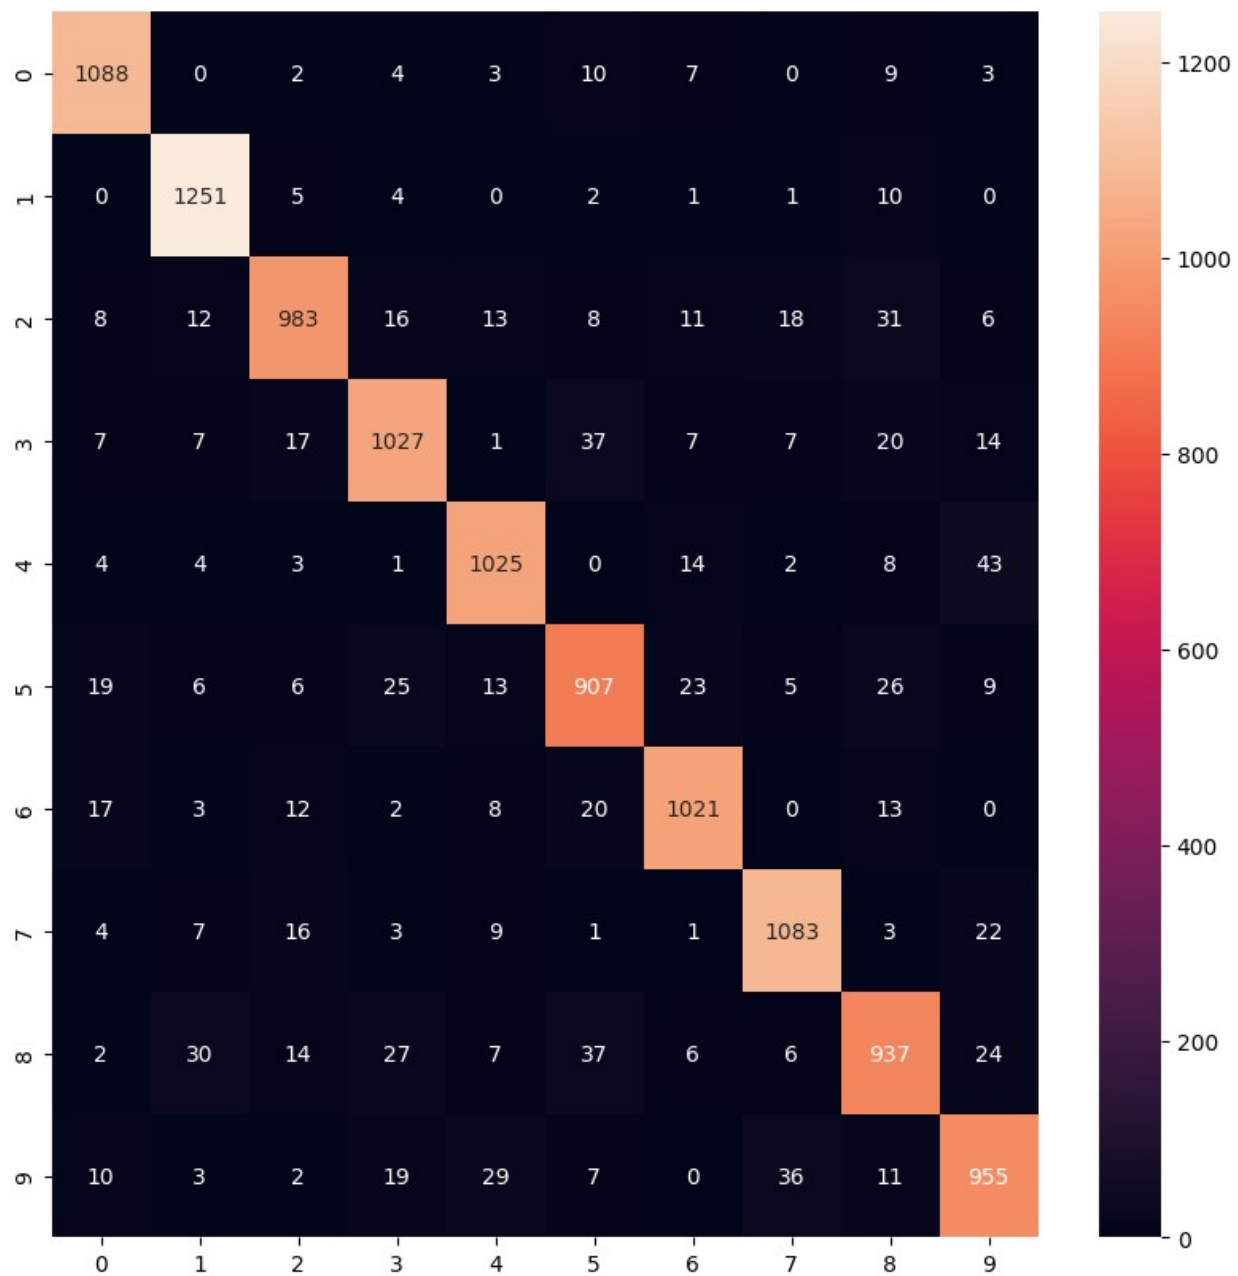
```
<IPython.core.display.HTML object>
```

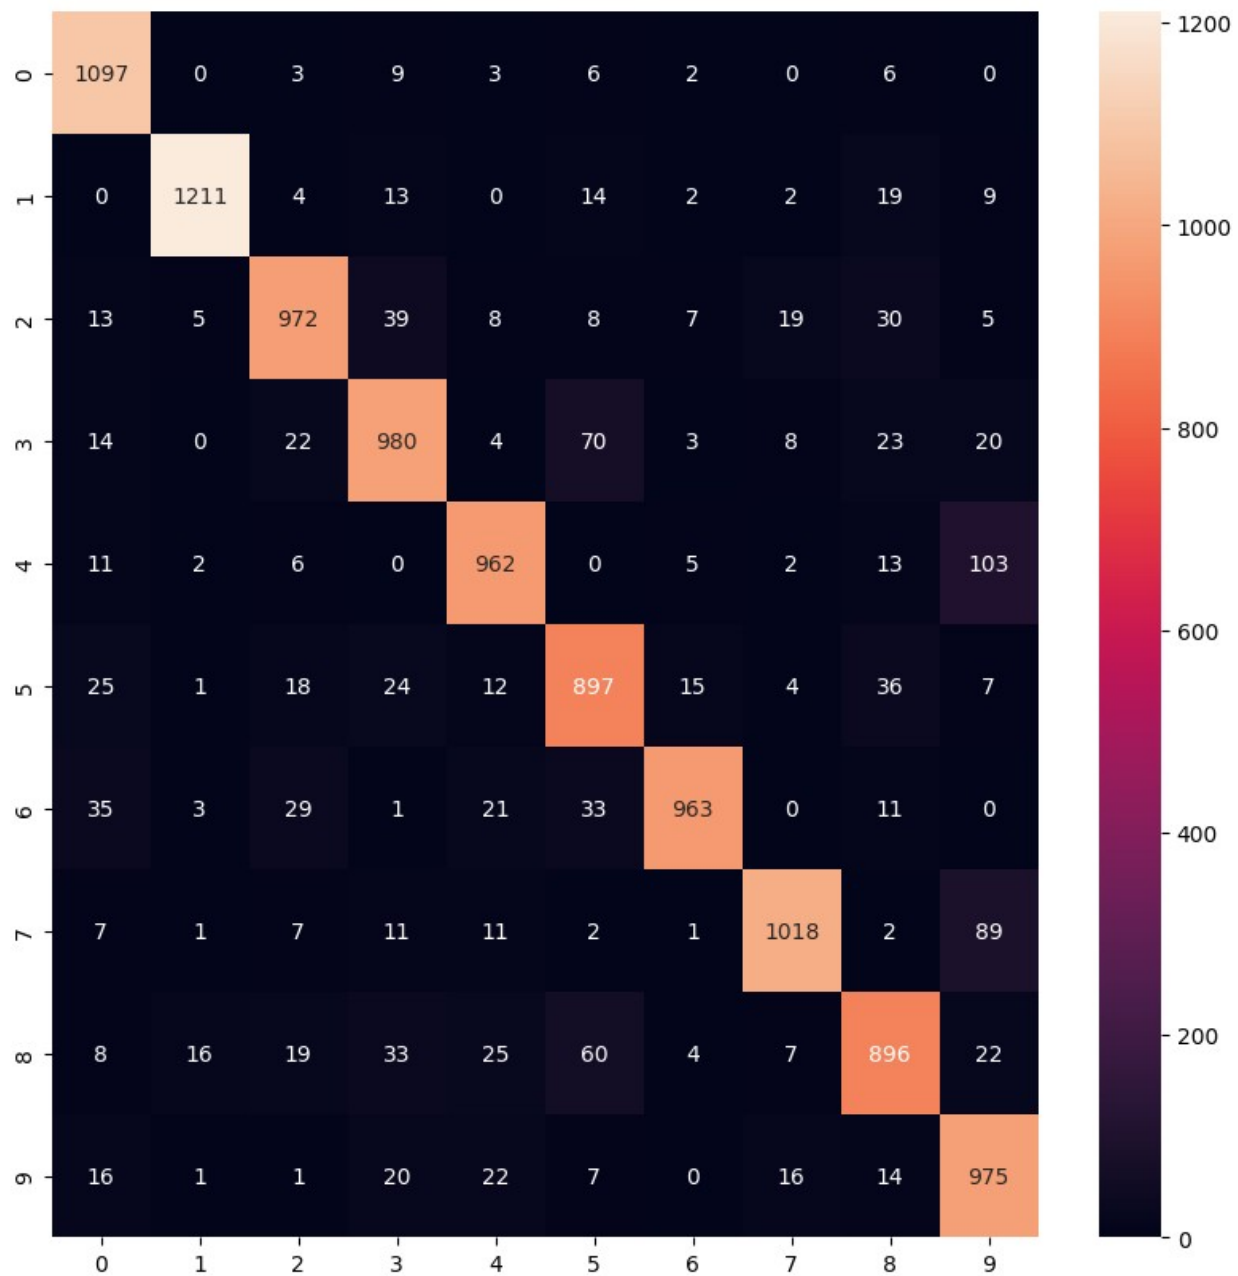
```
<IPython.core.display.HTML object>
```

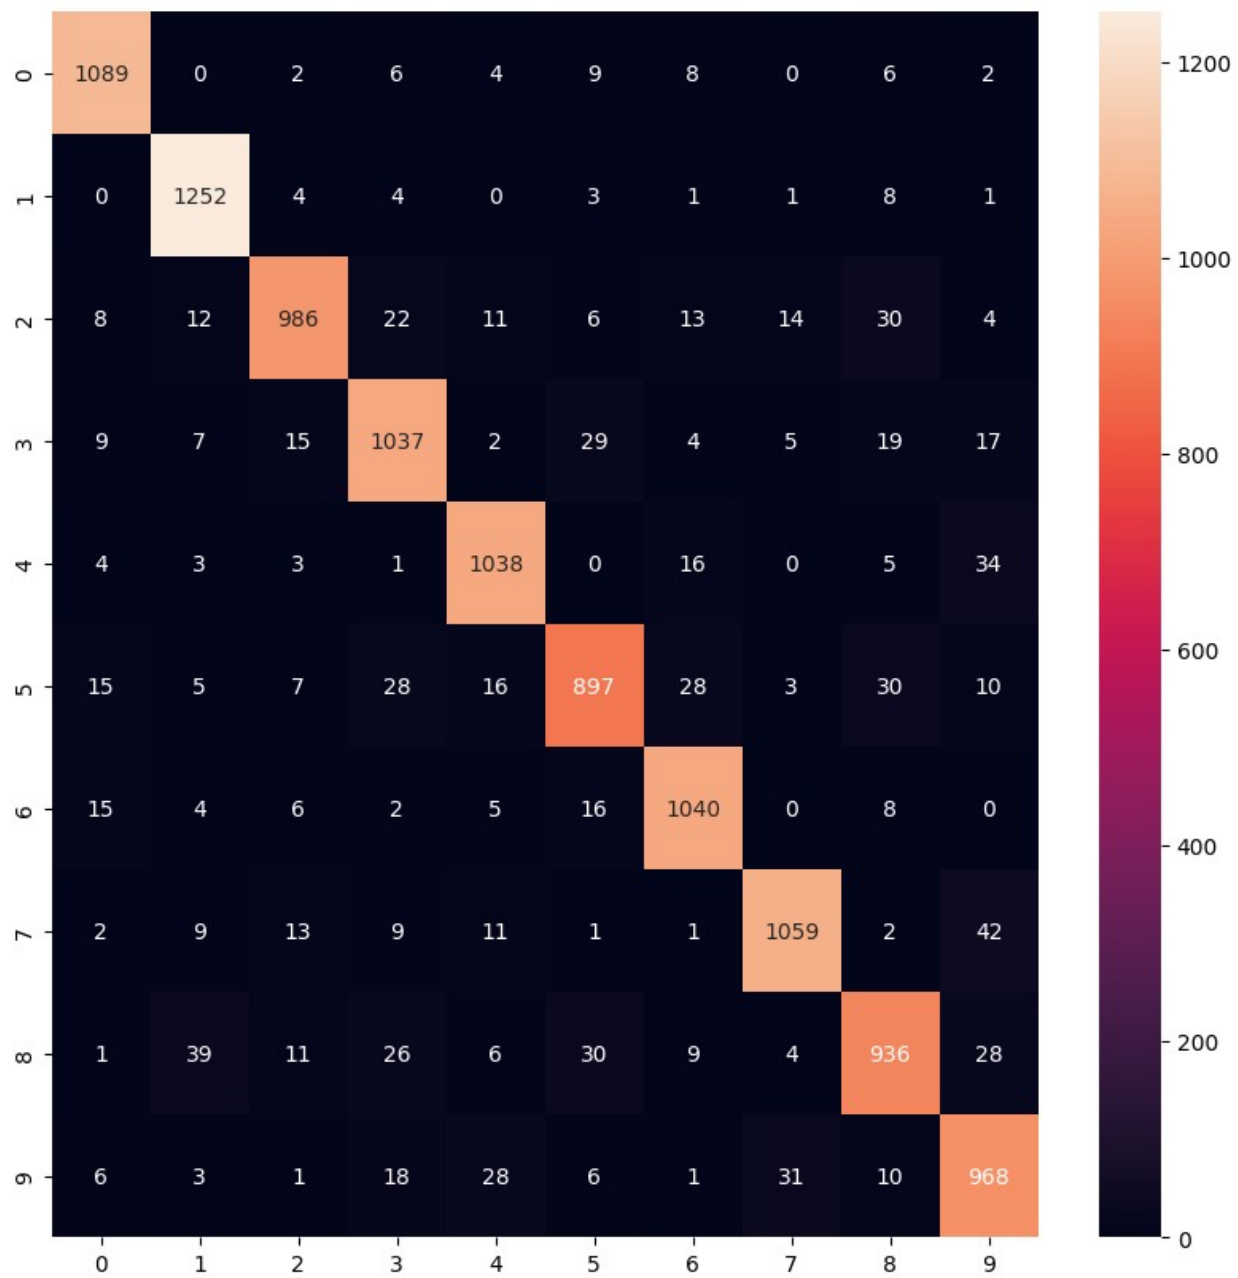



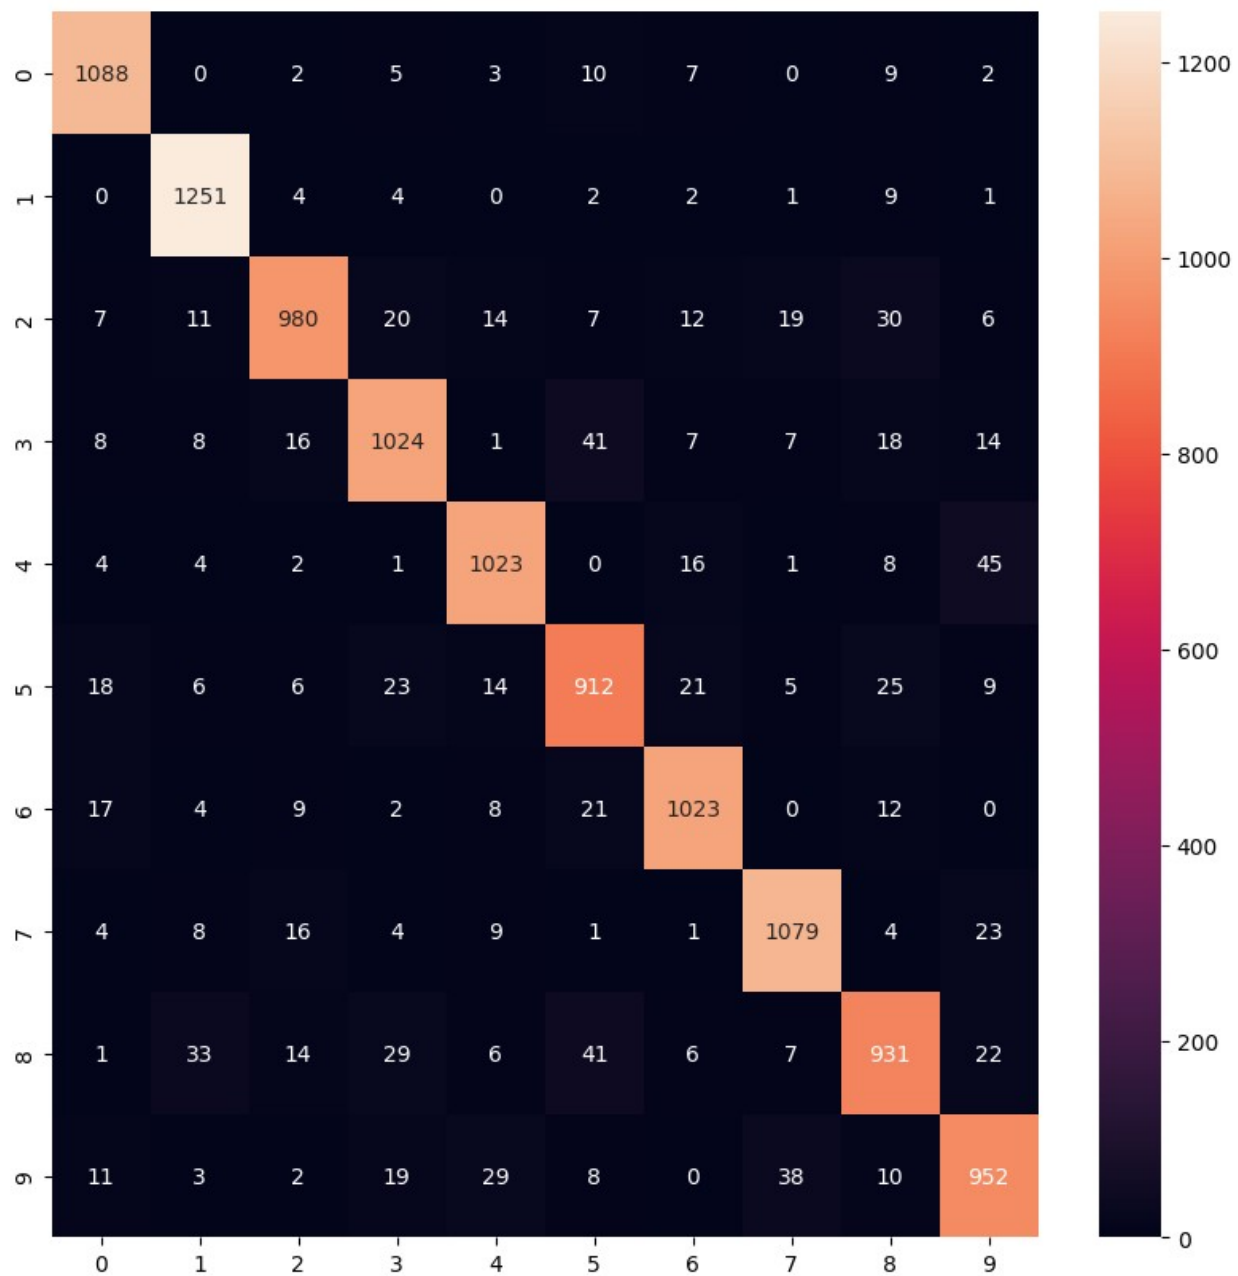


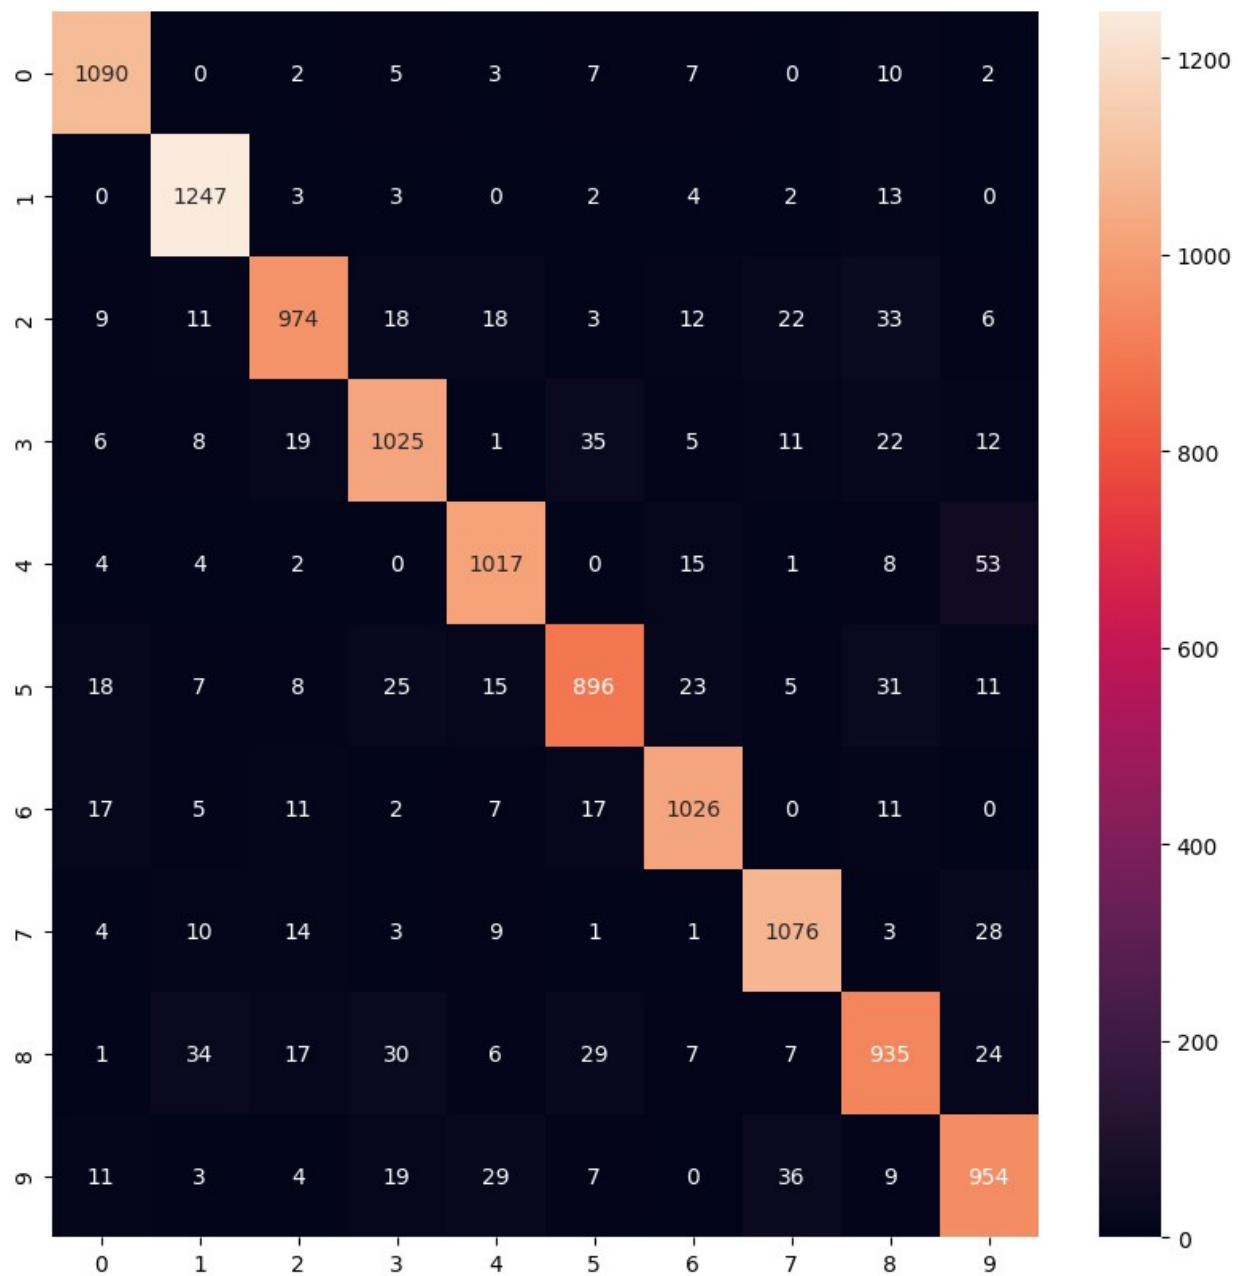


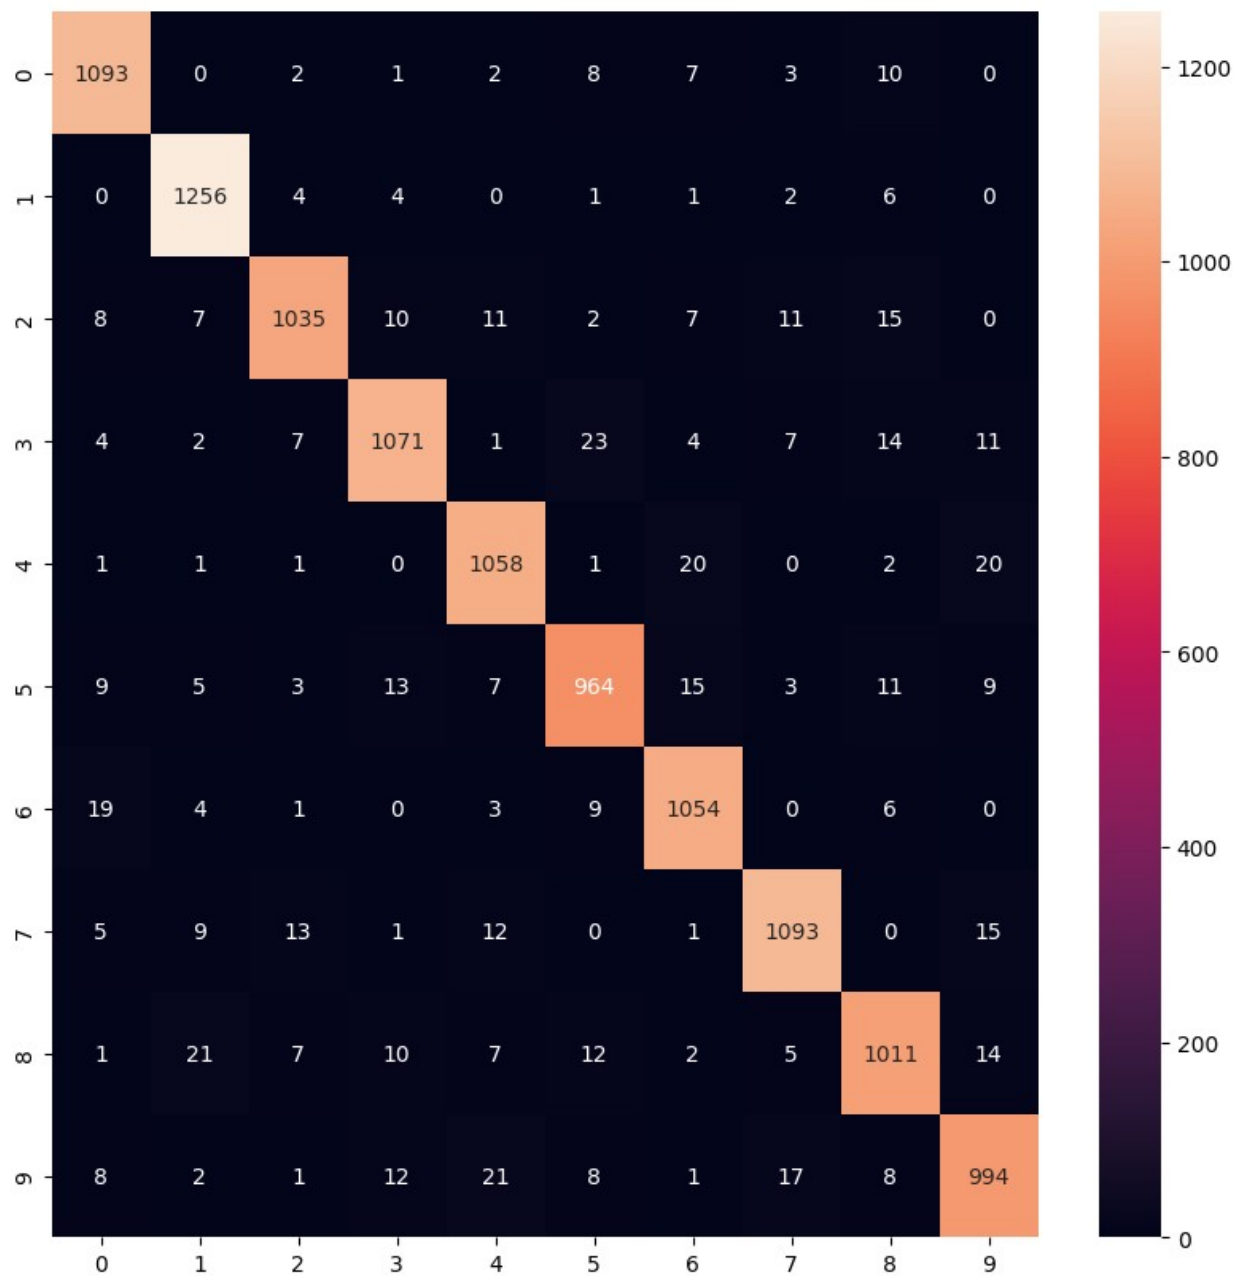


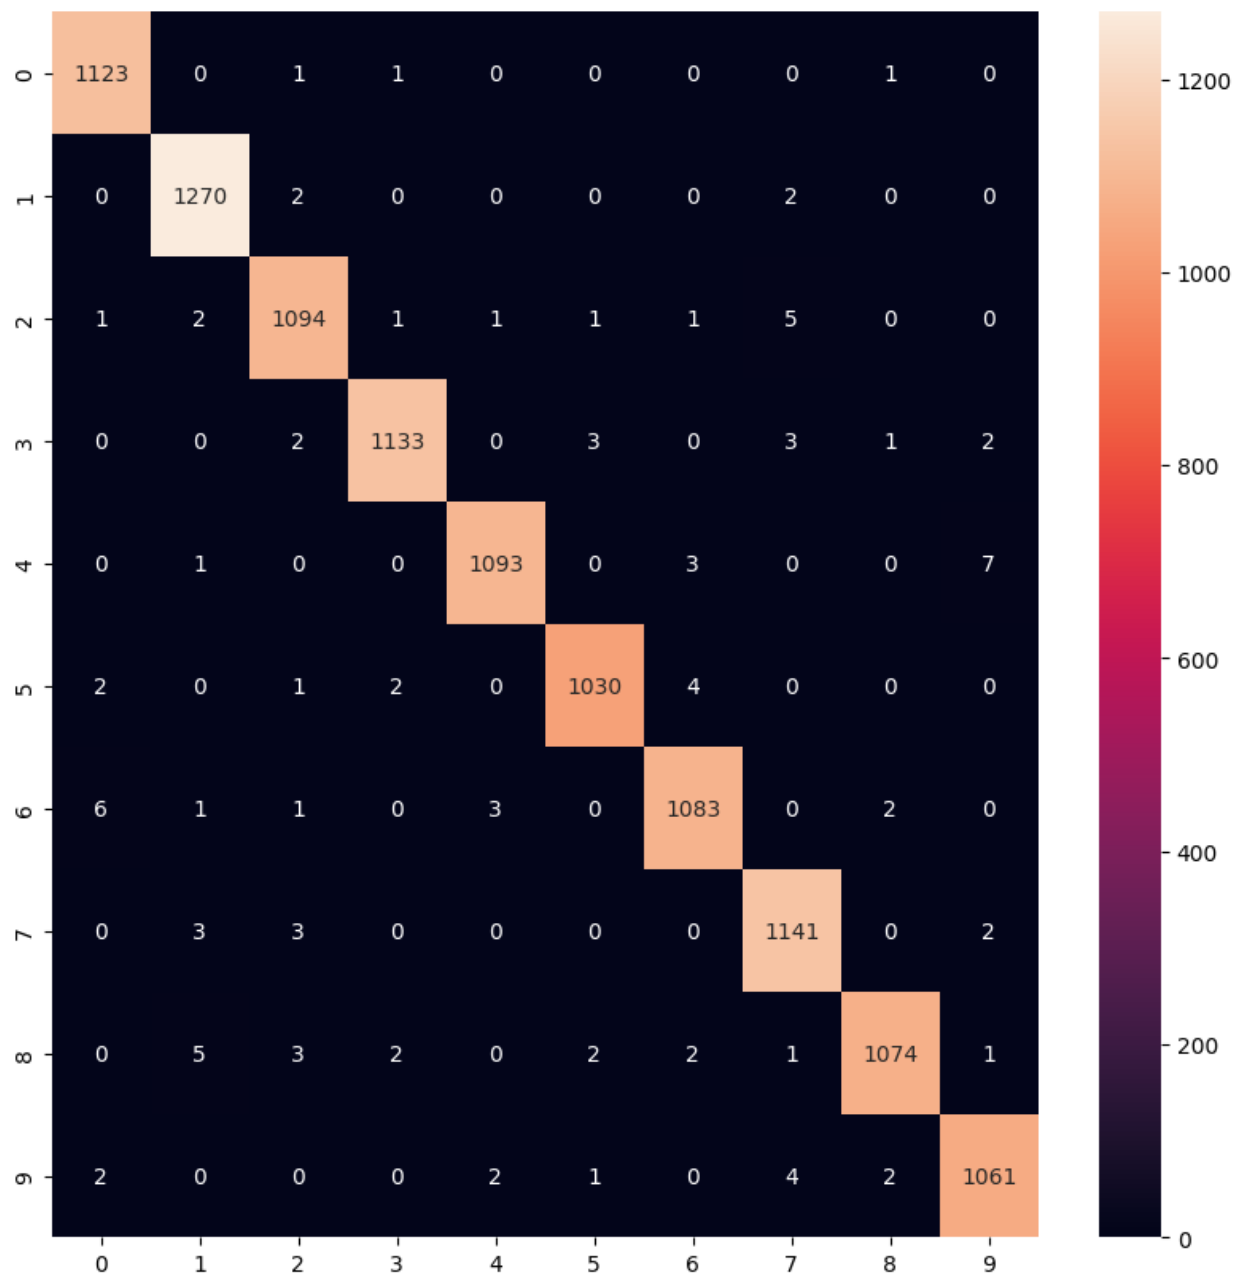


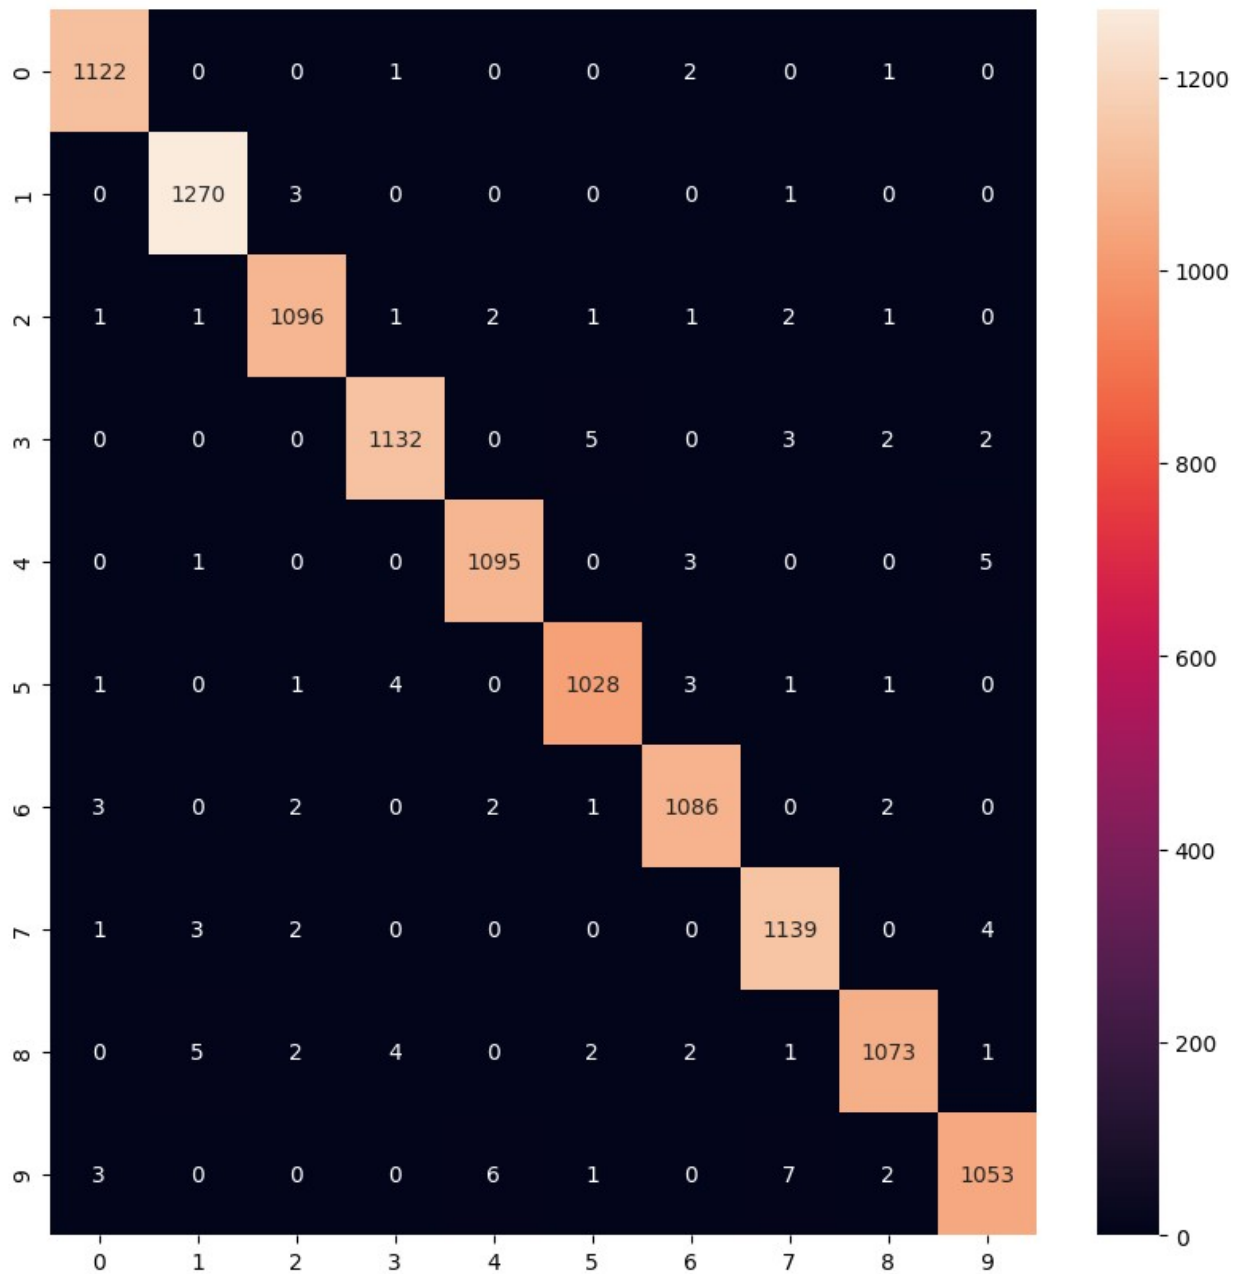












Evaluation

There are some evaluation insights in the notebook, but for a detailed evaluation, please refer to the [Wandb dashboard](#).

```
# Display summary table
summary_df = tracker.get_all_summaries()
# remove confusion matrix from summary
display(summary_df.drop(columns=['confusion_matrix']).style.highlight_max(subset=['accuracy']))
```

```
<pandas.io.formats.style.Styler at 0x7bdd5de2a4d0>
```

Model Comparison

Accuracy Analysis

- Best performing models:
 - a. `deep_net_medium`: 99.13% accuracy
 - b. `deep_net_large`: 99.05% accuracy
 - c. `deep_net_nano`: 94.90% accuracy
- Softmax regression models range between 71-92% accuracy

Learning Rate Impact

- `softmax_lr_0.01`: Best softmax regression model (92.14% accuracy)
- `softmax_lr_1.0`: Moderate performance (89.00% accuracy)
 - Indicates learning rate is too high
- `softmax_lr_0.00001`: Poor performance (71.54% accuracy)
 - Indicates learning rate is crucial for model convergence
 - Too low learning rate leads to slow convergence

Batch Size Experiments

- `softmax_batch_32`: Optimal batch size for softmax (91.98% accuracy)
- `softmax_batch_1`: Lowest performance (89.03% accuracy) with much longer training time.
- `softmax_batch_256`: Slightly lower performance (91.63% accuracy)

Deep Network Insights

- Deeper networks consistently outperform softmax regression
- `deep_net_medium` shows marginal improvement over `deep_net_large`
- Increasing network complexity doesn't always guarantee better performance

Inference Time Analysis

- Softmax models: ~0.0002-0.0003 seconds per inference
- Deep networks: ~0.0007-0.0015 seconds per inference
- Trade-off between accuracy and computational complexity

Training Time Observations

- Softmax models: 389-463 seconds
- Deep networks: 503-2887 seconds
- Deep networks require more training time due to increased complexity

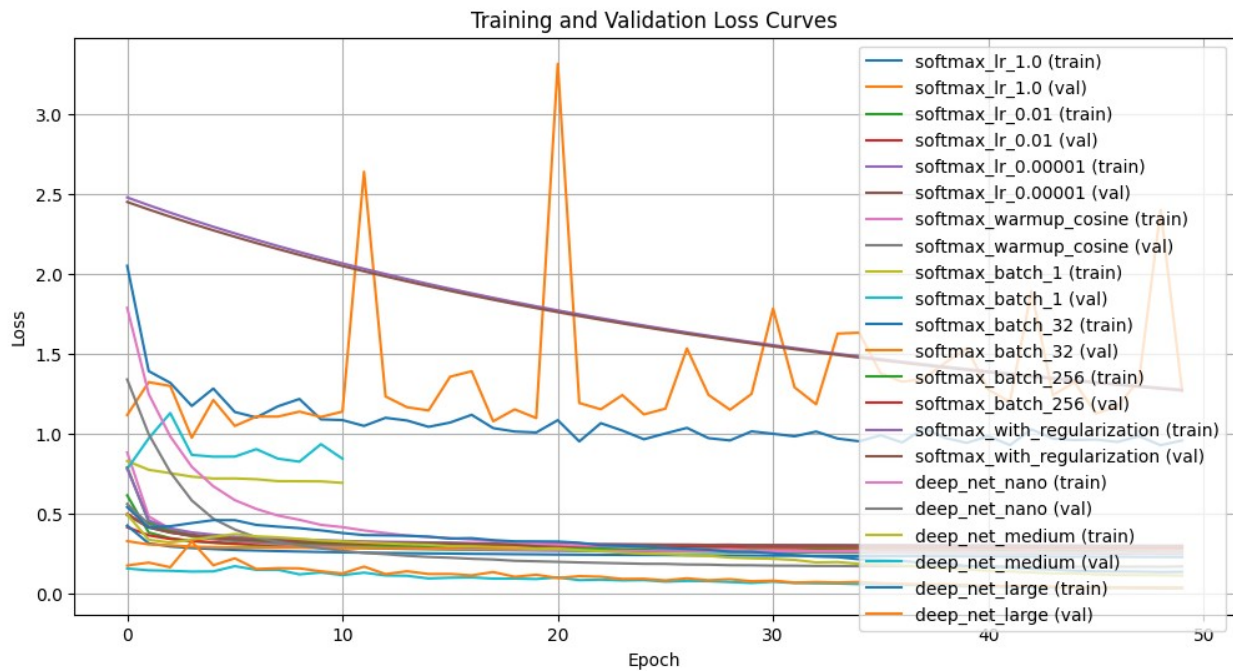
Regularization Impact

- `softmax_with_regularization`: Slight performance reduction
- Suggests careful hyperparameter tuning needed

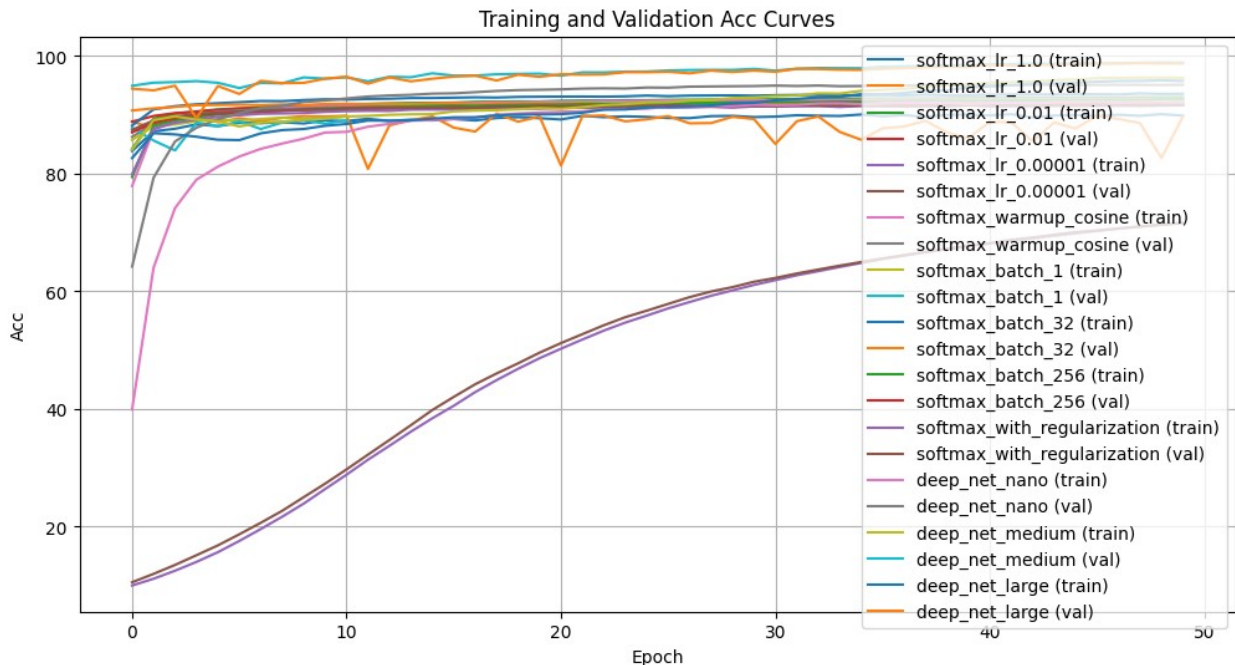
Key Takeaways

1. Deep networks superior for MNIST
2. Learning rate critical for model performance
3. Batch size influences training speed and accuracy
4. Computational complexity increases with network depth

```
# Plot training curves
fig = tracker.plot_training_curves('loss')
plt.show()
```



```
fig = tracker.plot_training_curves('acc')
plt.show()
```



Analysis of Training and Validation Curves

Key Observations

1. Softmax Models:

- `softmax_lr_1.0`:
 - Exhibits unstable training behavior with spikes in loss values.
 - Indicates that a high learning rate prevents smooth convergence.
- `softmax_lr_0.01`:
 - Best performance among softmax models, achieving the lowest validation loss.
 - Shows smoother training and validation curves, suggesting good learning rate.
- `softmax_lr_0.00001`:
 - High and slow-decreasing loss values throughout training.
 - Reflects that a learning rate too small leads to slow or incomplete convergence.
- Batch Size Impact:
 - Smaller batches (e.g., `softmax_batch_1`) result in noisier curves due to less accurate gradient estimates.
 - Larger batches (e.g., `softmax_batch_256`) stabilize the curves but sacrifice slight accuracy.

2. Deep Networks:

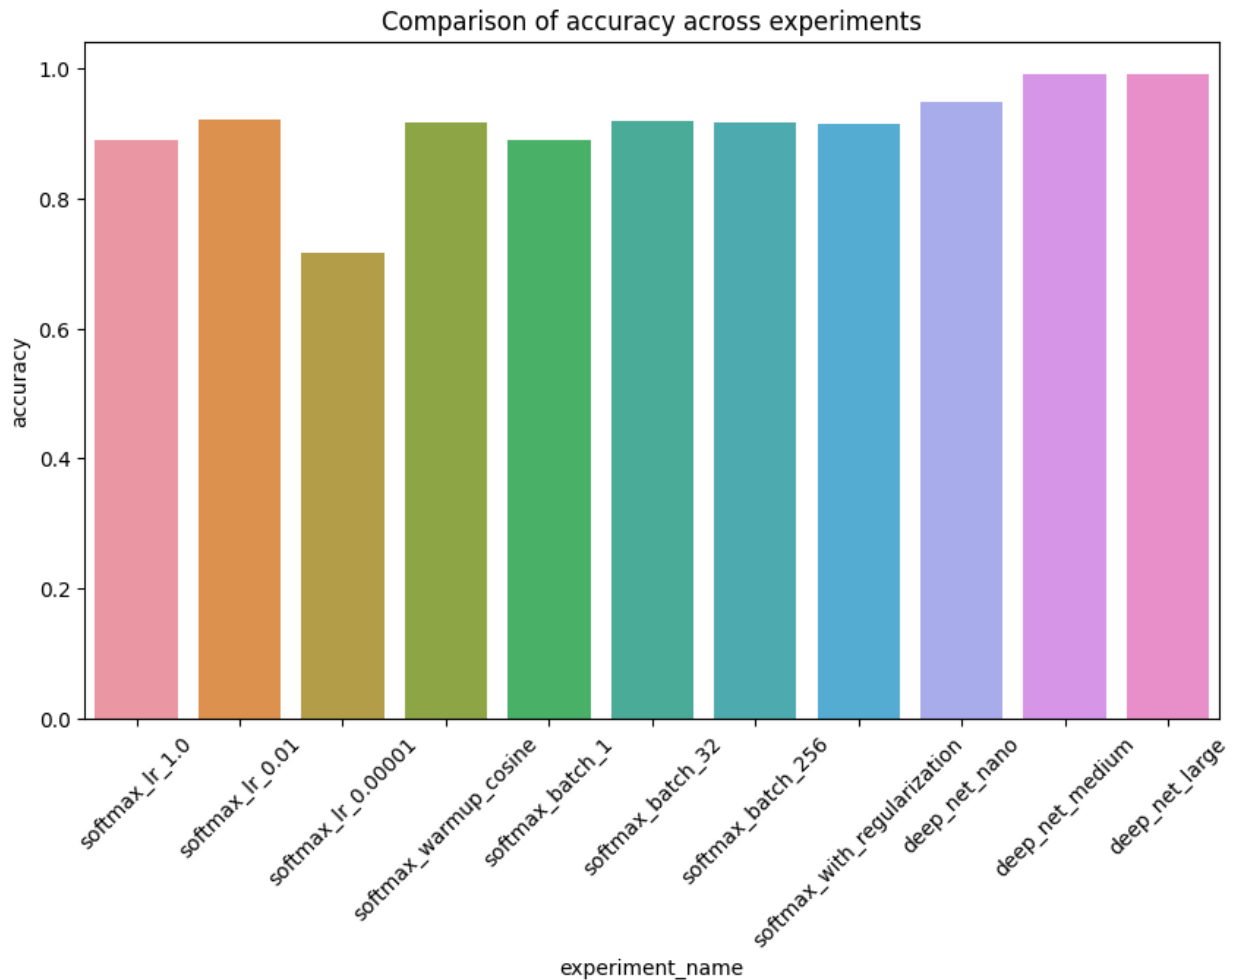
- Models like `deep_net_medium` and `deep_net_large` consistently show lower training and validation losses compared to softmax models.
- Network Size vs. Convergence:

- Smaller networks like `deep_net_nano` converge faster but reach lower accuracy values.
 - Larger networks take longer to converge but achieve better accuracy.
3. Regularization:
 - Regularization models show slightly elevated loss values compared to unregularized models, indicating the penalty imposed by regularization to prevent overfitting.
 4. Cosine Warmup Learning Rate:
 - The `softmax_warmup_cosine` curve demonstrates a gradual reduction in loss, indicating effective scheduling for stable convergence.

Insights from the Curves

1. Learning Rate:
 - Proper tuning is critical, as seen in the superior performance of `softmax_lr_0.01` compared to other learning rate configurations.
 - Extremely low or high learning rates lead to poor convergence and unstable training.
2. Model Complexity:
 - Deep networks outperform softmax models in terms of achieving lower loss and higher accuracy.
 - Increasing network complexity doesn't always yield significant performance improvement (e.g., `deep_net_medium` vs. `deep_net_large`).
3. Batch Size:
 - Optimal batch sizes balance stability and accuracy, as evidenced by the superior performance of `softmax_batch_32`.
4. Overfitting Indicators:
 - Narrow gaps between training and validation loss curves indicate better generalization, as seen in `deep_net_medium`.
5. Stability and Convergence:
 - Models with smoother loss curves (e.g., `deep_net_medium` and `softmax_lr_0.01`) are more reliable and generalize better than those with fluctuating loss patterns.

```
# Compare final metrics on test set
fig = tracker.plot_experiment_comparison('accuracy')
plt.show()
```



```
# Export results
summary_df.to_csv('experiment_results.csv')
```

```
# Detailed view of specific experiment
exp_df = tracker.get_experiment_df('deep_net_medium')
display(exp_df)
```

	epoch	train_loss	train_acc	val_loss	val_acc	
learning_rate \						
0	0	0.497291	84.109375	0.159649	94.942857	0.004667
1	1	0.338992	89.111607	0.148122	95.471429	0.006000
2	2	0.321408	89.743304	0.145222	95.571429	0.007333
3	3	0.340370	89.200893	0.140176	95.735714	0.008667
4	4	0.356799	88.883929	0.141500	95.450000	0.010000
5	5	0.374468	88.024554	0.174082	94.557143	0.009988

6	6	0.360194	88.705357	0.151320	95.435714	0.009951
7	7	0.354372	88.816964	0.151848	95.357143	0.009891
8	8	0.344741	89.189732	0.122903	96.364286	0.009807
9	9	0.334645	89.392857	0.134888	96.164286	0.009699
10	10	0.328464	89.738839	0.118059	96.371429	0.009568
11	11	0.322392	89.819196	0.133623	95.692857	0.009415
12	12	0.316809	90.008929	0.116012	96.471429	0.009241
13	13	0.311356	90.098214	0.114108	96.392857	0.009046
14	14	0.307336	90.133929	0.097346	97.064286	0.008831
15	15	0.298610	90.589286	0.101906	96.671429	0.008598
16	16	0.297787	90.720982	0.103948	96.678571	0.008347
17	17	0.293910	90.794643	0.096233	96.900000	0.008080
18	18	0.291423	90.948661	0.096302	96.964286	0.007798
19	19	0.282541	91.131696	0.093726	97.007143	0.007502
20	20	0.284431	91.033482	0.103723	96.657143	0.007195
21	21	0.274697	91.321429	0.086767	97.221429	0.006876
22	22	0.266992	91.658482	0.089871	97.214286	0.006549
23	23	0.258531	91.805804	0.084989	97.342857	0.006213
24	24	0.252172	92.109375	0.086919	97.364286	0.005872
25	25	0.248957	92.162946	0.077618	97.550000	0.005527
26	26	0.243560	92.412946	0.081955	97.621429	0.005179
27	27	0.238646	92.497768	0.080512	97.642857	0.004831
28	28	0.228051	92.863839	0.075118	97.642857	0.004483
29	29	0.223390	92.979911	0.069270	97.814286	0.004138
30	30	0.219947	93.198661	0.077390	97.514286	0.003797
31	31	0.213256	93.270089	0.068730	97.850000	0.003461

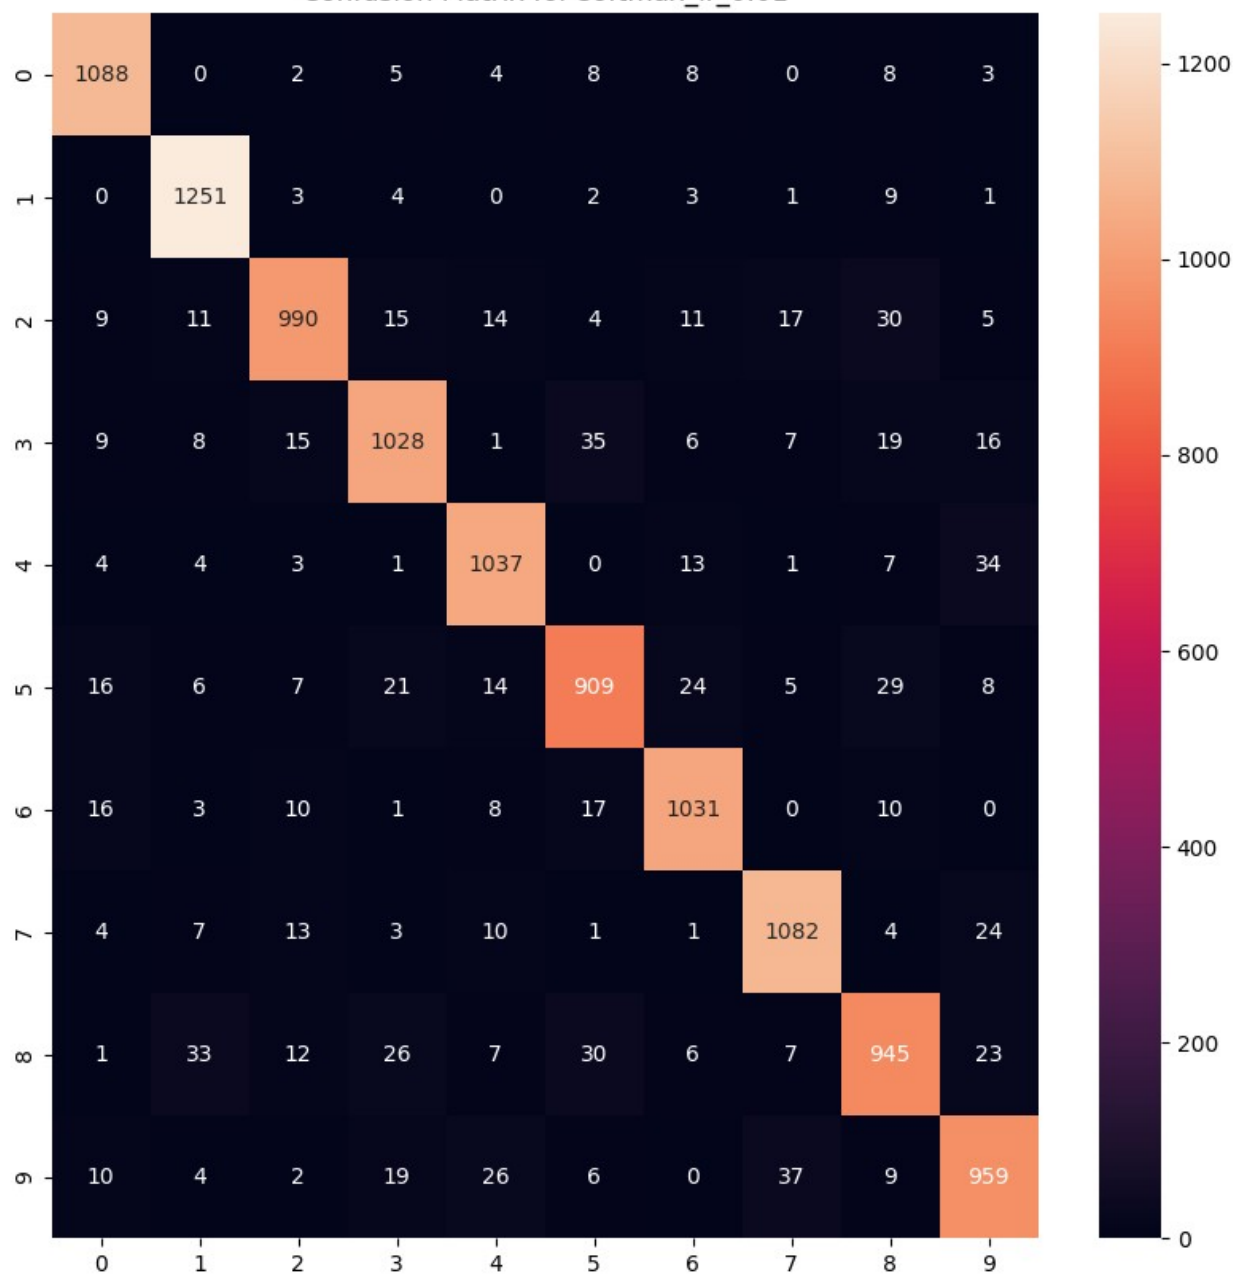
32	32	0.197659	93.662946	0.067091	97.964286	0.003134
33	33	0.199889	93.622768	0.065623	97.942857	0.002815
34	34	0.188858	94.082589	0.062283	97.928571	0.002508
35	35	0.184824	94.250000	0.058486	98.178571	0.002212
36	36	0.178929	94.415179	0.056118	98.242857	0.001930
37	37	0.169910	94.756696	0.055078	98.307143	0.001663
38	38	0.158018	95.064732	0.050002	98.442857	0.001412
39	39	0.153615	95.169643	0.047878	98.478571	0.001179
40	40	0.145442	95.466518	0.047226	98.557143	0.000964
41	41	0.141253	95.564732	0.043872	98.600000	0.000769
42	42	0.134877	95.665179	0.041059	98.728571	0.000595
43	43	0.128707	95.924107	0.040305	98.700000	0.000442
44	44	0.128374	95.988839	0.040662	98.735714	0.000311
45	45	0.121996	96.247768	0.039315	98.771429	0.000203
46	46	0.118953	96.287946	0.037799	98.771429	0.000119
47	47	0.117970	96.319196	0.036895	98.850000	0.000059
48	48	0.117851	96.350446	0.036186	98.878571	0.000022
49	49	0.115109	96.274554	0.036407	98.871429	0.000010

	epoch_time
0	55.996205
1	55.830900
2	55.088057
3	55.864079
4	54.722889
5	55.230920
6	55.298661
7	55.741643
8	55.441713
9	54.673023
10	55.307901
11	55.262882

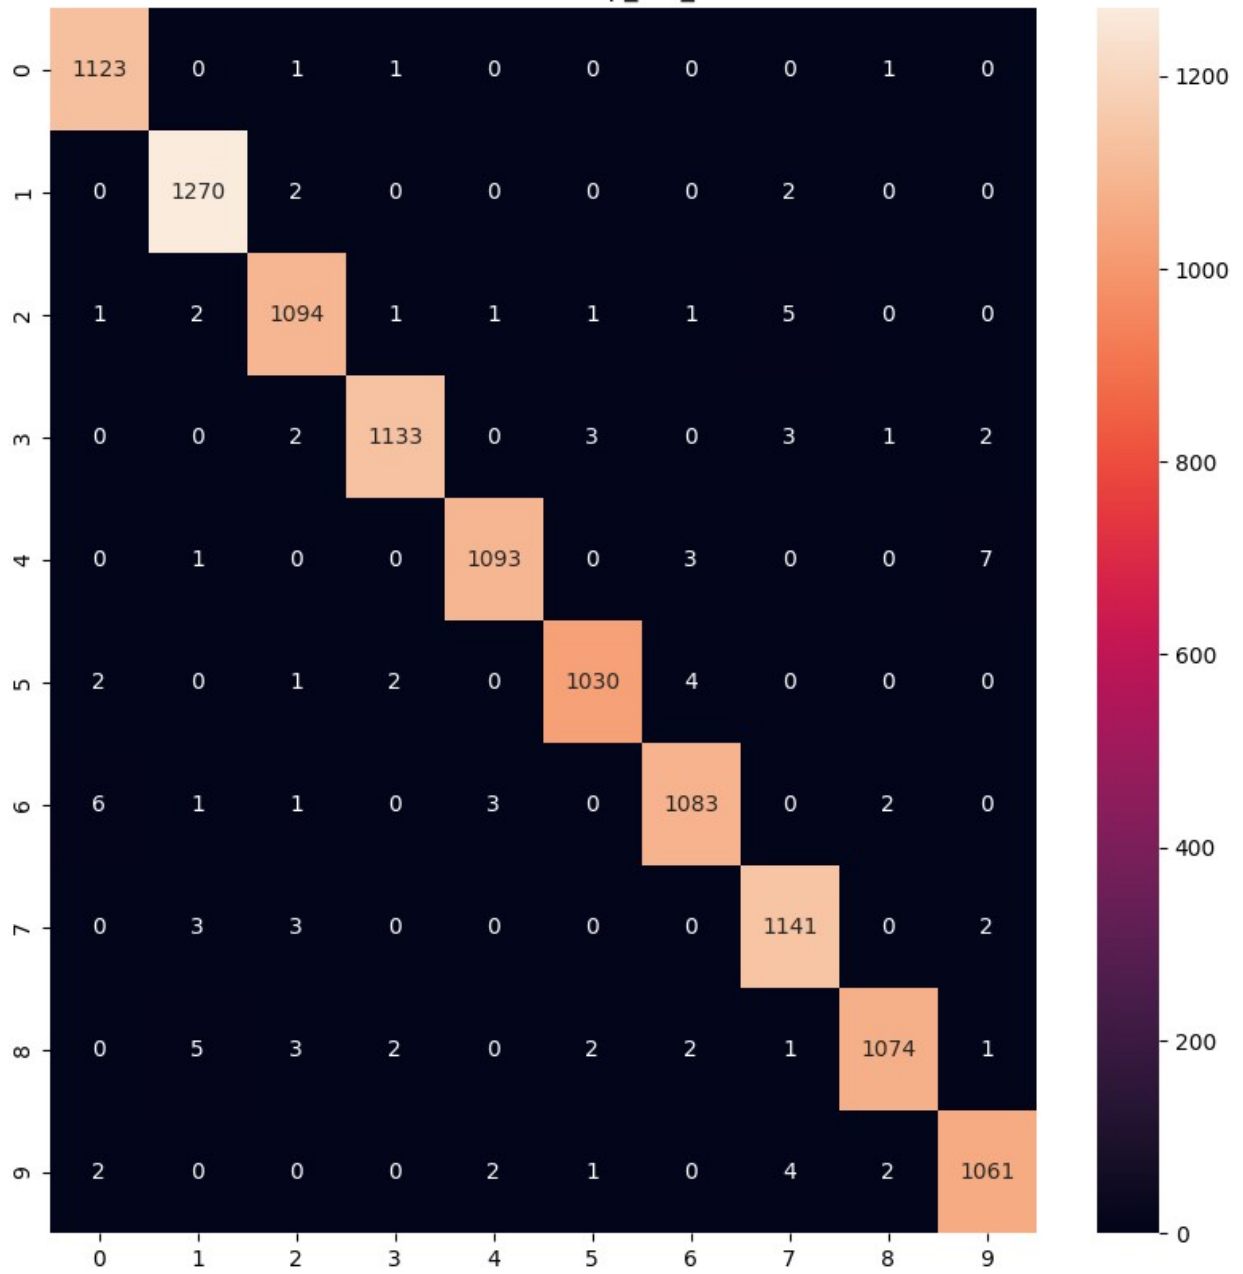
12	54.352537
13	56.216399
14	57.066223
15	56.571239
16	56.457256
17	55.229942
18	55.152944
19	55.387287
20	57.551653
21	56.896703
22	57.808463
23	55.984190
24	56.250254
25	56.417300
26	56.342187
27	56.036708
28	55.699501
29	56.235218
30	55.848260
31	55.152828
32	55.881460
33	55.737539
34	55.819153
35	55.711973
36	56.376323
37	56.225123
38	55.842007
39	55.884803
40	55.952831
41	55.170945
42	55.267149
43	56.017612
44	72.692998
45	56.677200
46	57.785577
47	56.317838
48	55.671846
49	56.070641

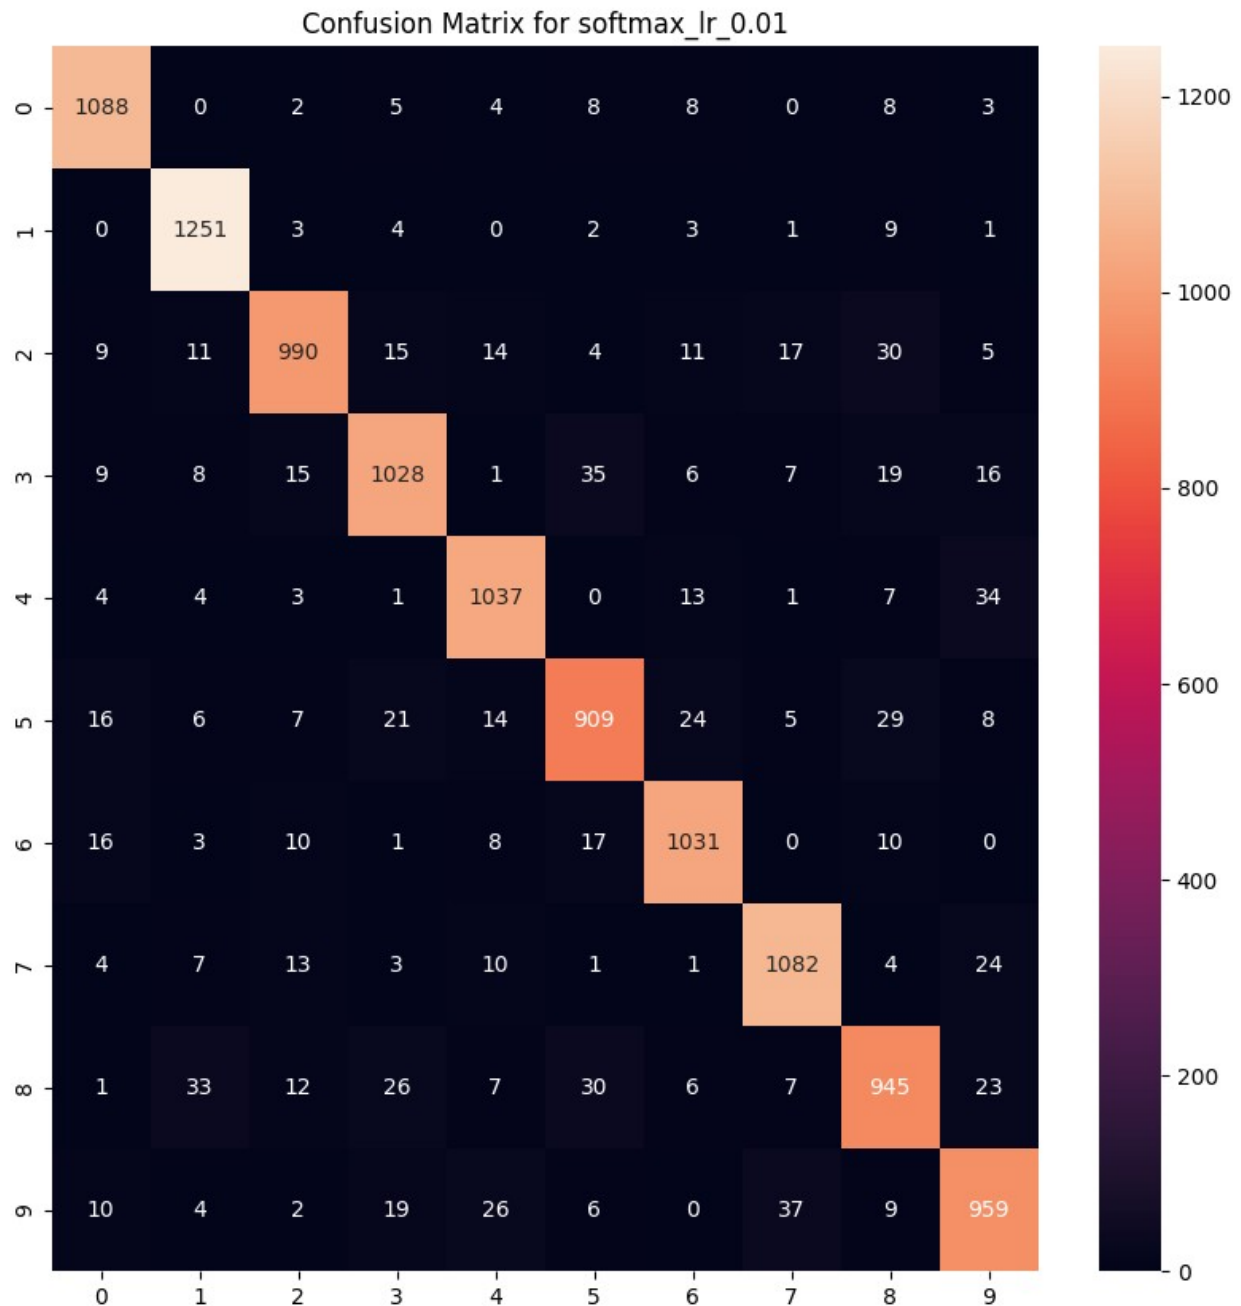
```
tracker.plot_confusion_matrix('deep_net_medium')  
tracker.plot_confusion_matrix('softmax_lr_0.01')
```

Confusion Matrix for softmax_lr_0.01



Confusion Matrix for deep_net_medium





Confusion Matrix Analysis

- The confusion matrix for the softmax regression model with the learning rate of 0.01 shows that the model is performing well, with relatively low misclassifications. However, it still struggles with certain digits, especially those that are visually similar e.g., 3vs5 or 4vs9.
- The confusion matrix of the deep neural network reveals an exceptionally high accuracy rate, with very few misclassifications. Most of the confusion happens between digits that are similar, like 4 vs 9, but these are minimal.

The deep neural network shows that it can effectively differentiate between a wide range of digit variations, which is reflected in the low false positive and false negative rates.

Misclassification Analysis on Deep Neural Network

We can see from the sample of the misclassified images that the deep neural network is struggling with digits that are not clearly written or are ambiguous. This is expected as even humans can struggle with these images.

